Context-Aware Notification Engine (CANE)

Deep Q-Network (DQN)

This notebook implements a **Deep Q-Network (DQN)** agent for the Context-Aware Notification Engine (CANE).

DQN models the notification pacing problem as a full Markov Decision Process (MDP). It approximates the optimal action-value function $Q^*(s, a)$ using a deep neural network, trained via Bellman error minimization with experience replay and a stationary target network:

$$\mathcal{L}(\theta) = \mathbb{E}_{(s,a,r,s',d)\sim\mathcal{D}}\left[\left(r + \gamma \cdot \max_{a'\in\mathcal{A}} Q_{\text{target}}(s', a') \cdot (1-d) - Q_{\text{online}}(s, a; \theta)\right)^2\right]$$

Notebook Structure

Section	Contents
1	Setup & Dependencies
2	Agent Interface
3	Environment Dynamics (`CANEEnv`)
4	Baselines & Evaluation Protocol
5	Environment Sanity Checks & Plotting Utilities
6	Deep Q-Network Architecture & Agent Implementation
7	Correctness Verification Tests
8	Registry & Execution Driver
9	Learning Curves & Optimisation Diagnostics
10	Decision Distribution & Policy Behaviour
11	Hyperparameter Sensitivity Analysis
12	Statistical Robustness & Model Persistence

---

1. Setup & Dependencies

Import required libraries, define action spaces, observation feature layouts, and initialize execution parameters.

In [ ]:
```python
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from abc import ABC, abstractmethod
import torch
import torch.nn as nn

# --- Action Space ------------------------------------------------------------
HOLD = 0        # Stay silent; fatigue decays
ENGAGE = 1      # Engagement notification
INCENTIVE = 2   # Incentive / promo notification

N_ACTIONS = 3
ACTION_NAMES = {HOLD: "Hold", ENGAGE: "Engage", INCENTIVE: "Incentive"}

# --- Observation Feature Indices ---------------------------------------------
IDX_HOUR = 0        # Raw hour (0-23)
IDX_DAY = 1         # Raw day of week (0-6)
IDX_ACTIVE = 2      # In-app activity flag (0 or 1)
IDX_FATIGUE = 3     # Accumulated fatigue [0, 1]
IDX_RECENCY = 4     # Normalised time since last send [0, 1]

N_BLOCKS = 8        # 3-hour blocks per day
BLOCK_HOURS = 24 // N_BLOCKS

IDX_ACT_RATE = 5    # Block activity rate (5..12)
IDX_CLICK_RATE = 13 # Block click rate (13..20)
IDX_TYPE_CR = 21    # Click rate per message type (21..22)
IDX_SINCE_CLICK = 23# Hours since last click

STATE_DIM_BASE = 5
STATE_DIM = 24

BELIEF_PRIOR = (1.0, 4.0)
USE_BELIEF = True

def agent_state_dim():
    return STATE_DIM if USE_BELIEF else STATE_DIM_BASE

DEVICE = "cpu"
torch.set_num_threads(4)

print("Actions:", ACTION_NAMES)
print("State dim:", STATE_DIM, "| agent sees:", agent_state_dim())
print("torch version:", torch.__version__, "| device:", DEVICE)
```

```
Actions: {0: 'Hold', 1: 'Engage', 2: 'Incentive'}
State dim: 24 | agent sees: 24
torch version: 2.6.0+cu124 | device: cpu
```

---

2. Agent Interface

Abstract base class defining the shared execution contract for all agents.

In [18]:
```python
# --- The agent contract ------------------------------------------------
# Every policy in this notebook -- the two non-learning baselines and the
# learner alike -- implements this interface, which is what lets the evaluation
# harness below treat them all identically. The comparison is only like-for-like
# because no agent gets a code path of its own.
#
# `act` returns (action, aux). The aux dict is how an agent hands the harness
# whatever it needs to record: the deep agents put their Q-values in it.
# `update` is a no-op for the baselines; `reset` is called at every episode
# boundary, which matters for any agent holding per-episode state.
class Agent(ABC):
    """Abstract base class for all CANE agents."""

    @property
    def name(self):
        return self.__class__.__name__

    @abstractmethod
    def act(self, state, greedy=False):
        """Select an action given an observation state."""
        raise NotImplementedError

    @abstractmethod
    def update(self, state, action, reward, next_state, done, aux):
        """Update policy / value function from a transition."""
        raise NotImplementedError

    def reset(self):
        """Reset agent state at episode boundaries (default no-op)."""
        pass

print("Agent interface defined.")
```

```
Agent interface defined.
```

---

3. Environment Dynamics

Defines user archetypes, click/fatigue/churn dynamics, and the `CANEEnv` environment.

In [19]:
```python
# --- Environment Parameters --------------------------------------------
ARCHETYPE_W = {
    "OfficeWorker":     {ENGAGE: 1.15, INCENTIVE: 1.05},
    "NightOwlStudent":  {ENGAGE: 0.90, INCENTIVE: 1.60},
    "NightShiftWorker": {ENGAGE: 1.00, INCENTIVE: 1.25},
    "NormalStudent":    {ENGAGE: 0.90, INCENTIVE: 1.55},
    "Housewife":        {ENGAGE: 1.05, INCENTIVE: 1.45},
}

CANE_CONFIG = dict(
    steps_per_episode=168,
    lam=0.90,
    kappa={HOLD: 0.0, ENGAGE: 0.12, INCENTIVE: 0.18},
```

```
    mu=0.80,
    w_arch=ARCHETYPE_W,
    churn_threshold=0.70,
    gamma_0=-7.0,
    gamma_1=4.0,
    R_click=10.0,
    W_send={HOLD: 0.0, ENGAGE: 1.0, INCENTIVE: 2.5},
    W_fat=2.0,
    R_churn=30.0,
    beta=1.15,
    streak_mode="marginal",
    streak_cap=5.0,
    streak_lapse_hours=36,
    active_scale=0.6,
)

ARCHETYPES = ["OfficeWorker", "NightOwlStudent", "NightShiftWorker",
              "NormalStudent", "Housewife"]

def _bump(hours, centre, width, amp):
    delta = (hours - centre + 12.0) % 24.0 - 12.0
    return amp * np.exp(-0.5 * (delta / width) ** 2)

def archetype_curve(archetype, hours=None):
    h = np.arange(24.0) if hours is None else np.asarray(hours, dtype=float)
    if archetype == "OfficeWorker":
        c = _bump(h, 7.5, 1.0, 0.42) + _bump(h, 12.5, 1.1, 0.30) + _bump(h, 19.5, 2
    elif archetype == "NightOwlStudent":
        c = _bump(h, 0.0, 2.2, 0.32) + _bump(h, 22.0, 1.6, 0.26) + _bump(h, 16.0, 2
    elif archetype == "NightShiftWorker":
        c = _bump(h, 16.5, 1.6, 0.45) + _bump(h, 6.5, 1.5, 0.40) + _bump(h, 2.0, 1.
    elif archetype == "NormalStudent":
        c = _bump(h, 16.0, 1.6, 0.45) + _bump(h, 21.0, 1.5, 0.42) + _bump(h, 7.0, 1
    elif archetype == "Housewife":
        c = _bump(h, 10.0, 2.4, 0.42) + _bump(h, 14.5, 2.4, 0.40) + _bump(h, 20.5,
    else:
        raise ValueError(f"unknown archetype: {archetype}")
    return np.clip(c, 0.0, 1.0)

ARCHETYPE_TABLE = {a: archetype_curve(a) for a in ARCHETYPES}
WEEKEND_FLATTEN = {"OfficeWorker": 0.55, "NormalStudent": 0.50,
                   "NightShiftWorker": 0.30, "NightOwlStudent": 0.15,
                   "Housewife": 0.05}

print(f"{len(ARCHETYPES)} archetypes defined:", ", ".join(ARCHETYPES))
```

```
5 archetypes defined: OfficeWorker, NightOwlStudent, NightShiftWorker, NormalStuden
t, Housewife
```

In [20]:

```
# --- The simulated user -------------------------------------------------
# The environment is the assignment's model of one app user over a week: 168
# hourly steps, each of which is a decision point. There is no external dataset;
# this class IS the data-generating process, and the archetype tables above are
# its ground truth.
#
# The agent never observes the archetype or the true receptiveness curve. It
```

```python
# sees only the belief state assembled in `_observe`, so learning who this
# person is has to happen through reward alone. That is the whole problem.
class CANEEnv:
    """Context-Aware Notification Engine simulator environment."""

    def __init__(self, config=None, archetype=None, seed=0):
        self.cfg = {**CANE_CONFIG, **(config or {})}
        self.fixed_archetype = archetype
        self.rng = np.random.default_rng(seed)
        self.reset()

    def _p0(self):
        # Baseline receptiveness for this person at this hour of the day.
        # On weekends (day >= 5) it is blended toward the person's daily mean:
        # a weekday commuter's sharp 08:00 peak flattens out on a Saturday.
        base = ARCHETYPE_TABLE[self.archetype][self.hour]
        if self.day >= 5:
            f = WEEKEND_FLATTEN[self.archetype]
            base = (1.0 - f) * base + f * ARCHETYPE_TABLE[self.archetype].mean()
        return float(base)

    def _observe(self):
        # The belief state. Note what is absent: the archetype, the true
        # receptiveness curve, and the click probability. The agent gets only
        # what an app could actually log about a user.
        s = np.zeros(STATE_DIM, dtype=np.float32)
        s[IDX_HOUR] = self.hour
        s[IDX_DAY] = self.day
        s[IDX_ACTIVE] = self.active
        s[IDX_FATIGUE] = self.fatigue
        s[IDX_RECENCY] = self.recency
        # Beta-prior smoothing. Early in an episode a block may have zero sends,
        # and a raw clicks/sends ratio would be 0/0. Adding the prior makes an
        # unobserved block read as 'unknown' rather than as 'known to be dead'.
        a0, b0 = BELIEF_PRIOR
        s[IDX_ACT_RATE:IDX_ACT_RATE + N_BLOCKS] = (self.blk_act_hits + a0) / (self.
        s[IDX_CLICK_RATE:IDX_CLICK_RATE + N_BLOCKS] = (self.blk_clicks + a0) / (sel
        s[IDX_TYPE_CR] = (self.type_clicks[ENGAGE] + a0) / (self.type_sends[ENGAGE]
        s[IDX_TYPE_CR + 1] = (self.type_clicks[INCENTIVE] + a0) / (self.type_sends[
        s[IDX_SINCE_CLICK] = min(self.hours_since_click / 48.0, 1.0)
        return s

    def _roll_activity(self):
        # Whether the user is already in the app this hour. Rolled every step,
        # including on Hold, because the app observes app-opens whether or not it
        # sent anything -- that is the free signal the belief state accumulates.
        p = min(1.0, self.cfg["active_scale"] * self._p0())
        self.active = int(self.rng.random() < p)
        b = self.hour // BLOCK_HOURS
        self.blk_act_obs[b] += 1
        self.blk_act_hits[b] += self.active

    def reset(self, seed=None):
        if seed is not None:
            self.rng = np.random.default_rng(seed)
        # Episodes start at a random hour and day, so the agent cannot learn a
```

```
        # fixed step-index-to-hour mapping and must read the hour from the state.
        self.archetype = self.fixed_archetype or ARCHETYPES[self.rng.integers(len(A
        self.t = 0
        self.hour = int(self.rng.integers(24))
        self.day = int(self.rng.integers(7))
        self.fatigue = 0.0
        self.recency = 1.0
        self.hours_since_send = 24
        self.streak = 0
        self.hours_since_click = 0
        self.opted_out = False
        self.blk_act_obs = np.zeros(N_BLOCKS)
        self.blk_act_hits = np.zeros(N_BLOCKS)
        self.blk_sends = np.zeros(N_BLOCKS)
        self.blk_clicks = np.zeros(N_BLOCKS)
        self.type_sends = np.zeros(N_ACTIONS)
        self.type_clicks = np.zeros(N_ACTIONS)
        self._roll_activity()
        return self._observe(), {"archetype": self.archetype}

    def step(self, action):
        cfg = self.cfg
        action = int(action)
        sent = action != HOLD
        f_now = self.fatigue

        clicked = False
        # A click is only possible on a send, and only when the user is NOT
        # already active in the app: a notification delivered to someone already
        # using the app earns nothing. Click probability then scales the hour's
        # receptiveness down by current fatigue and by the per-archetype weight
        # of this message type.
        if sent and not self.active:
            w_a = cfg["w_arch"][self.archetype][action]
            p_click = np.clip(self._p0() * (1.0 - cfg["mu"] * f_now) * w_a, 0.0, 1.
            clicked = bool(self.rng.random() < p_click)

        blk = self.hour // BLOCK_HOURS
        if sent:
            self.blk_sends[blk] += 1
            self.blk_clicks[blk] += clicked
            self.type_sends[action] += 1
            self.type_clicks[action] += clicked

        if clicked:
            self.streak += 1
            self.hours_since_click = 0
        else:
            self.hours_since_click += 1
            if self.hours_since_click > cfg["streak_lapse_hours"]:
                self.streak = 0

        # Consecutive clicks compound, which is what makes this a sequential
        # problem rather than a per-hour one: the value of a click depends on
        # whether the previous ones landed.
        streak_bonus = 0.0
```

```
        if clicked and self.streak > 0:
            if cfg["streak_mode"] == "cumulative":
                streak_bonus = sum(cfg["beta"] ** k for k in range(1, self.streak +
            else:
                streak_bonus = cfg["beta"] ** self.streak
            if cfg["streak_cap"] is not None:
                streak_bonus = min(streak_bonus, cfg["streak_cap"])

        # Fatigue decays geometrically at `lam` and takes on `kappa[action]` per
        # send. Because lam < 1, one send's cost is paid back over many future
        # steps -- this recursion is the reason a send is expensive and the
        # reason a patient agent can rationally choose to send nothing at all.
        self.fatigue = float(np.clip(cfg["lam"] * f_now + cfg["kappa"][action], 0.0

        # Past the threshold the user may opt out permanently, ending the
        # episode. This is the hard ceiling on over-sending: no recovery.
        if self.fatigue > cfg["churn_threshold"]:
            hazard = 1.0 / (1.0 + np.exp(-(cfg["gamma_0"] + cfg["gamma_1"] * self.f
            self.opted_out = bool(self.rng.random() < hazard)

        # Reward = click payoff, minus the cost of sending, minus a penalty on
        # fatigue, minus the churn charge, plus any streak bonus.
        # Fatigue is charged at `f_now`, the level BEFORE this step's send, so a
        # send is not billed for the fatigue it has only just created -- that
        # arrives on the following steps via the recursion above.
        reward = (cfg["R_click"] * clicked
                  - cfg["W_send"][action]
                  - cfg["W_fat"] * f_now
                  - cfg["R_churn"] * self.opted_out
                  + streak_bonus)

        self.t += 1
        self.hour = (self.hour + 1) % 24
        if self.hour == 0:
            self.day = (self.day + 1) % 7

        # Recency, capped at 24h and normalised, is the agent's only direct view
        # of how recently it last made contact.
        self.hours_since_send = 0 if sent else self.hours_since_send + 1
        self.recency = float(min(self.hours_since_send / 24.0, 1.0))
        self._roll_activity()

        terminated = self.opted_out
        truncated = self.t >= cfg["steps_per_episode"]

        info = {
            "archetype": self.archetype,
            "clicked": clicked, "sent": sent,
            "fatigue": self.fatigue, "fatigue_pre": f_now,
            "opted_out": self.opted_out, "streak": self.streak,
            "streak_bonus": streak_bonus,
            "hour": int((self.hour - 1) % 24), "active_pre": self.active,
        }
        return self._observe(), float(reward), terminated, truncated, info

print("CANEEnv defined.")
```

```
CANEEnv defined.
```

---

4. Baselines & Evaluation Protocol

Provides benchmark agents (`FixedScheduleAgent` , `RandomAgent`) and evaluation harness functions.

In [ ]:
```python
# --- Non-learning baselines ---------------------------------------------
# Fixed-18:00 is the industry default this project argues against: one
# hard-coded send time for everyone. Random is the lower bound. Neither
# learns, so both are evaluated once rather than across training seeds.
class FixedScheduleAgent(Agent):
    """Sends an Engagement Nudge at a fixed hour every day."""
    def __init__(self, hour=18, action=ENGAGE):
        self.hour, self.action = hour, action

    @property
    def name(self):
        return f"Fixed-{self.hour:02d}:00"

    def act(self, state, greedy=False):
        return (self.action if int(state[IDX_HOUR]) == self.hour else HOLD), {}

    def update(self, *args, **kwargs):
        return {}


class RandomAgent(Agent):
    """Selects actions uniformly at random."""
    def __init__(self, seed=0):
        self.rng = np.random.default_rng(seed)

    @property
    def name(self):
        return "Random"

    def act(self, state, greedy=False):
        return int(self.rng.integers(N_ACTIONS)), {}

    def update(self, *args, **kwargs):
        return {}


# --- The evaluation harness ---------------------------------------------
# One function drives both training and evaluation, which is what keeps
# them honest: `learn=True, greedy=False` trains, `learn=False,
# greedy=True` scores a frozen policy. Passing explicit `seeds` replays
# an identical set of episodes for every agent.
def run_episodes(agent, env, n_episodes=None, seeds=None, learn=True,
                 greedy=False, collect_log=False):
    """Runs agent in env over specified episodes and returns evaluation metrics."""
    if seeds is None:
```

```
        seeds = [None] * int(n_episodes)

    ep_rewards, ep_sends, ep_clicks, ep_optouts, ep_lengths = [], [], [], [], []
    # Counts actions by clock hour, which is what the policy heatmaps and the
    # personalisation analysis are built from later.
    hour_actions = np.zeros((24, N_ACTIONS), dtype=int)
    log = []

    for ep, sd in enumerate(seeds):
        state, info = env.reset(seed=sd)
        agent.reset()
        total_r = sends = clicks = 0.0

        while True:
            action, aux = agent.act(state, greedy=greedy)
            next_state, reward, terminated, truncated, info = env.step(action)

            if learn:
                agent.update(state, action, reward, next_state, terminated, aux)

            total_r += reward
            sends += info["sent"]
            clicks += info["clicked"]
            hour_actions[info["hour"], action] += 1

            if collect_log:
                log.append({"episode": ep, "step": env.t, "hour": info["hour"],
                            "archetype": info["archetype"], "action": action,
                            "clicked": info["clicked"], "sent": info["sent"],
                            "fatigue": info["fatigue"], "reward": reward,
                            "streak": info["streak"], "opted_out": info["opted_out"

            state = next_state
            if terminated or truncated:
                ep_optouts.append(float(terminated))
                break

        ep_rewards.append(total_r)
        ep_sends.append(sends)
        ep_clicks.append(clicks)
        ep_lengths.append(env.t)

    # CTR is clicks per SEND, not per step, and is defined as 0.0 when the
    # agent never sent. A never-send policy therefore reports ctr 0.00 and
    # reward 0.00 -- neither is a missing value; both are what silence earns.
    total_sends = float(np.sum(ep_sends))
    metrics = {
        "agent": agent.name,
        "reward_mean": float(np.mean(ep_rewards)),
        "reward_std": float(np.std(ep_rewards)),
        "ctr": float(np.sum(ep_clicks) / total_sends) if total_sends > 0 else 0.0,
        "sends_per_episode": float(np.mean(ep_sends)),
        "clicks_per_episode": float(np.mean(ep_clicks)),
        "optout_rate": float(np.mean(ep_optouts)),
        "episode_length": float(np.mean(ep_lengths)),
    }
```

```
    return metrics, (log if collect_log else None), np.array(ep_rewards), hour_acti


# --- The shared evaluation protocol ------------------------------------
# These four values are identical in all four notebooks. That identity is
# what makes the cross-algorithm comparison valid, and `check_results.py`
# verifies it afterwards by confirming the deterministic baseline rows
# agree exactly across the four result files.
#
# The 900,000+ evaluation seeds are disjoint from the 1000+ training and
# 7000+ evaluation environment seeds below, so nothing is scored on an
# episode it was trained on. QUICK is a smoke-test switch: its numbers
# are NOT reportable.
QUICK = False

EVAL_SEEDS = list(range(900_000, 900_050 if QUICK else 900_200))
TRAIN_EPISODES = 150 if QUICK else 600
N_SEEDS = 2 if QUICK else 5

def preseed_buffer(agent, env, steps=1000):
    """Pre-populates the replay buffer with peak-hour sends so DQN starts with posi
    s, _ = env.reset()
    for _ in range(steps):
        hour = int(s[IDX_HOUR])
        # Smart exploration during typical peak activity windows (morning, lunch, e
        if hour in [7, 8, 12, 13, 18, 19, 20] and np.random.rand() < 0.4:
            a = ENGAGE
        else:
            a = HOLD if np.random.rand() < 0.7 else int(np.random.randint(N_ACTIONS

        ns, r, term, trunc, _ = env.step(a)
        agent.buffer.push(agent._features(s), a, r, agent._features(ns), float(term
        s = env.reset()[0] if (term or trunc) else ns


print(f"Eval protocol: {len(EVAL_SEEDS)} reserved episodes; "
      f"training budget {TRAIN_EPISODES} episodes; {N_SEEDS} seeds."
      + (" [QUICK]" if QUICK else ""))
```

```
Eval protocol: 200 reserved episodes; training budget 600 episodes; 5 seeds.
```

In [ ]:

```
# --- Frozen-policy probes ----------------------------------------------
# Snapshots taken during training, used for the learning curves and the
# animation. A separate seed block again, so probing never touches the
# episodes used for training or for the reported score.
PROBE_SEEDS = list(range(950_000, 950_020))
SNAPSHOT_ARCHETYPES = ("OfficeWorker", "NightOwlStudent")

def policy_snapshot(agent, archetypes=SNAPSHOT_ARCHETYPES, seeds=PROBE_SEEDS):
    """Evaluate frozen greedy policy across probe seeds."""
    out = {}
    for arch in archetypes:
        env = CANEEnv(seed=8000, archetype=arch)
        m, _, _, hours = run_episodes(agent, env, seeds=seeds, learn=False, greedy=
        out[arch] = {"hour_actions": hours, "reward": m["reward_mean"],
                     "ctr": m["ctr"], "sends": m["sends_per_episode"],
```

```
                    "optout": m["optout_rate"]}
    return out

# Trains a fresh agent per seed and scores each on the held-out episodes.
# A fresh agent per seed, not one agent re-evaluated: the spread being
# reported is the variance of the LEARNING process, not of the scoring.
def train_and_evaluate(make_agent, n_seeds=None, train_episodes=TRAIN_EPISODES,
                       archetype=None, collect_hours=False,
                       snapshot_every=None, snapshot_seed=0,
                       snapshot_archetypes=SNAPSHOT_ARCHETYPES):
    """Train an agent across multiple random seeds and evaluate on held-out seeds."
    n_seeds = N_SEEDS if n_seeds is None else n_seeds
    per_seed, hour_stack, snapshots = [], [], []

    for seed in range(n_seeds):
        agent = make_agent(seed)
        # Distinct environment seeds for training and evaluation, so a policy
        # cannot succeed by memorising particular episodes.
        train_env = CANEEnv(seed=1000 + seed, archetype=archetype)

        if hasattr(agent, "buffer"):
           preseed_buffer(agent, train_env, steps=1000)

        if snapshot_every and seed == snapshot_seed:
            snapshots.append((0, policy_snapshot(agent, snapshot_archetypes)))
            done = 0
            while done < train_episodes:
                chunk = min(snapshot_every, train_episodes - done)
                run_episodes(agent, train_env, n_episodes=chunk, learn=True, greedy
                done += chunk
                snapshots.append((done, policy_snapshot(agent, snapshot_archetypes)
        else:
            run_episodes(agent, train_env, n_episodes=train_episodes, learn=True, g

        eval_env = CANEEnv(seed=7000 + seed, archetype=archetype)
        m, _, rewards, hours = run_episodes(agent, eval_env, seeds=EVAL_SEEDS, lear
        per_seed.append(m)
        hour_stack.append(hours)

    rewards = np.array([m["reward_mean"] for m in per_seed])
    summary = {
        "agent": per_seed[0]["agent"],
        "reward_mean": rewards.mean(),
        "reward_std": rewards.std(ddof=0),
        # ddof=1: the sample standard deviation across seeds, which is what the
        # confidence intervals in the figures are derived from.
        "reward_sem_std": rewards.std(ddof=1) if n_seeds > 1 else 0.0,
        "ctr": np.mean([m["ctr"] for m in per_seed]),
        "sends_per_episode": np.mean([m["sends_per_episode"] for m in per_seed]),
        "optout_rate": np.mean([m["optout_rate"] for m in per_seed]),
        "per_seed_rewards": rewards,
    }
    if collect_hours:
        summary["hour_actions"] = np.sum(hour_stack, axis=0)
    if snapshots:
        summary["snapshots"] = snapshots
```

```
    return summary

# Baselines have nothing to train, so they are evaluated once. Both are
# deterministic given the seed list, which is precisely why their rows
# can be used to prove the four notebooks share one protocol.
def evaluate_baseline(make_agent, archetype=None):
    agent = make_agent(0)
    env = CANEEnv(seed=7000, archetype=archetype)
    m, _, rewards, hours = run_episodes(agent, env, seeds=EVAL_SEEDS, learn=False,
    return {"agent": m["agent"], "reward_mean": m["reward_mean"],
            "reward_std": m["reward_std"], "reward_sem_std": 0.0,
            "ctr": m["ctr"], "sends_per_episode": m["sends_per_episode"],
            "optout_rate": m["optout_rate"],
            "per_seed_rewards": np.array([m["reward_mean"]]),
            "hour_actions": hours}
```

---

5. Environment Sanity Checks & Exploratory Data Analysis (EDA)

Validates environment dynamics and visualises baseline responsiveness curves, fatigue dynamics, click surfaces, and ground-truth policy reference maps.

In [24]:
```
# --- 5.1 Environment Sanity Checks ------------------------------------------
env = CANEEnv(seed=0)

e = CANEEnv(archetype="Housewife", seed=1)
e.reset(seed=1)
f_trace, expected = [], 0.0
for _ in range(30):
    _, _, term, trunc, info = e.step(ENGAGE)
    expected = min(0.9 * expected + CANE_CONFIG["kappa"][ENGAGE], 1.0)
    f_trace.append((info["fatigue"], expected))
    if term or trunc:
        break
assert all(abs(a - b) < 1e-9 for a, b in f_trace)

e.reset(seed=2)
for _ in range(5):
    e.step(ENGAGE)
f_high = e.fatigue
for _ in range(10):
    e.step(HOLD)
assert e.fatigue < f_high

class _AlwaysHold(Agent):
    def act(self, s, greedy=False): return HOLD, {}
    def update(self, *a, **k): return {}

m_hold, _, _, _ = run_episodes(_AlwaysHold(), CANEEnv(seed=3), 30, learn=False, gre
assert m_hold["optout_rate"] == 0.0
```

```
class _AlwaysSend(Agent):
    def act(self, s, greedy=False): return ENGAGE, {}
    def update(self, *a, **k): return {}

m_spam, _, _, _ = run_episodes(_AlwaysSend(), CANEEnv(seed=4), 30, learn=False, gre
assert m_spam["optout_rate"] > 0.8

print("Environment sanity checks passed.")
```

```
Environment sanity checks passed.
```

In [25]:
```
# --- Plot G1: Receptive Windows by Archetype ---
from pathlib import Path

FOCUS = ["OfficeWorker", "NightOwlStudent"]
ARCH_COLOUR = {"OfficeWorker": "#2f6fb5", "NightOwlStudent": "#c4453c"}

FIGDIR = Path("figures")
FIGDIR.mkdir(exist_ok=True)

def save(fig, tag):
    fig.savefig(FIGDIR / f"{tag}.png", dpi=150, bbox_inches="tight")

def weekend_curve(archetype):
    base = ARCHETYPE_TABLE[archetype]
    f = WEEKEND_FLATTEN[archetype]
    return (1.0 - f) * base + f * base.mean()

hours = np.arange(24)
fig, axes = plt.subplots(1, 2, figsize=(10.5, 3.4), sharey=True)
for ax, arch in zip(axes, FOCUS):
    ax.fill_between(hours, ARCHETYPE_TABLE[arch], color=ARCH_COLOUR[arch], alpha=0.
    ax.plot(hours, ARCHETYPE_TABLE[arch], color=ARCH_COLOUR[arch], lw=2.0, label="w
    ax.plot(hours, weekend_curve(arch), color=ARCH_COLOUR[arch], lw=1.3, ls="--", l
    ax.set_title(arch, fontsize=10)
    ax.set_xlabel("hour of day")
    ax.set_xticks(range(0, 24, 3))
    ax.legend(frameon=False, fontsize=8)
axes[0].set_ylabel(r"$p_0$  click propensity")
fig.suptitle("G1   Receptive windows by archetype", fontsize=11)
fig.tight_layout()
save(fig, "G1_archetype_curves")
plt.show()
```

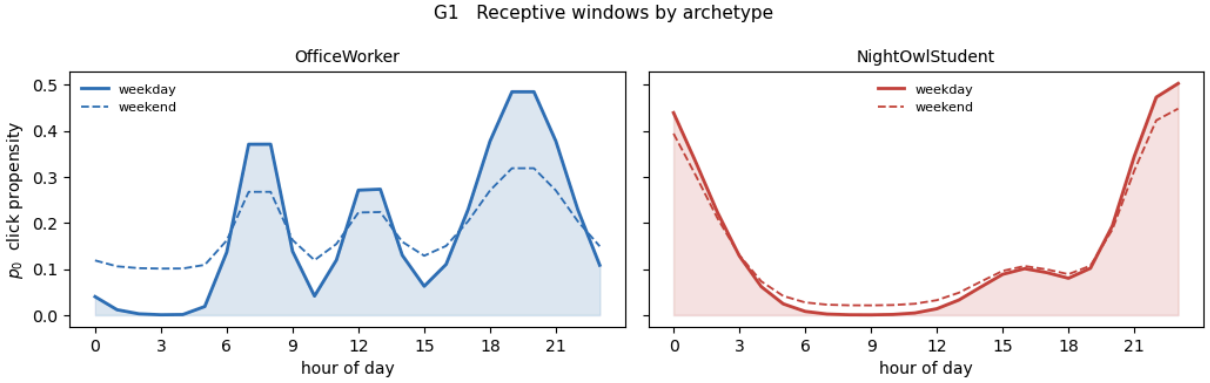

In [26]:
```
# --- Plot G2: Fatigue Dynamics ---
lam = CANE_CONFIG["lam"]
kap = CANE_CONFIG["kappa"]
thr = CANE_CONFIG["churn_threshold"]
pattern = np.array([1, 1, 1, 1, 1, 1] + [0] * 18)

fig, ax = plt.subplots(figsize=(9, 3.0))
ax.fill_between(range(len(pattern)), 0, 1, where=pattern.astype(bool),
                color="#999999", alpha=0.12, step="post")
for act, col in [(ENGAGE, "#2f6fb5"), (INCENTIVE, "#c4453c")]:
    F, trace = 0.0, []
    for a in pattern:
        F = float(np.clip(lam * F + (kap[act] if a else 0.0), 0.0, 1.0))
        trace.append(F)
    ax.step(range(len(trace)), trace, where="post", lw=1.8, color=col,
            label=f"{ACTION_NAMES[act]}  (kappa={kap[act]})")

ax.axhline(thr, color="#666666", ls=":", lw=1.4)
ax.text(len(pattern) - 0.5, thr + 0.015, f"churn threshold {thr}",
        ha="right", va="bottom", fontsize=8, color="#666666")
ax.set_xlabel("hours")
ax.set_ylabel("fatigue $F$")
ax.set_ylim(0, 1)
ax.legend(frameon=False, fontsize=8, loc="upper right")
ax.set_title("G2   Six hourly sends (shaded), then silence", fontsize=11)
fig.tight_layout()
save(fig, "G2_fatigue_dynamics")
plt.show()
```

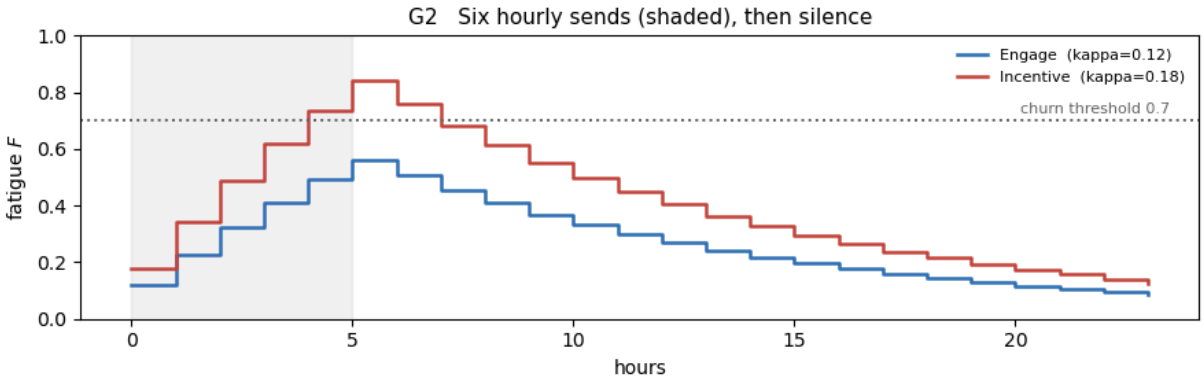

In [27]:
```
# --- Plot G3: Click Probability Surface ---
mu = CANE_CONFIG["mu"]

def break_even(archetype, action):
    cfg = CANE_CONFIG
    debt = cfg["W_fat"] * cfg["kappa"][action] / (1.0 - cfg["lam"])
    return (debt + cfg["W_send"][action]) / (cfg["R_click"] * cfg["w_arch"][archety

F_grid = np.linspace(0.0, 1.0, 101)
fig, axes = plt.subplots(1, 2, figsize=(11, 3.5), sharey=True)
for ax, arch in zip(axes, FOCUS):
    P = np.clip(np.outer(1.0 - mu * F_grid, ARCHETYPE_TABLE[arch])
                * CANE_CONFIG["w_arch"][arch][ENGAGE], 0.0, 1.0)
    im = ax.imshow(P, origin="lower", aspect="auto", cmap="magma",
```

```
                  extent=[0, 23, 0, 1], vmin=0, vmax=0.6)
    cs = ax.contour(np.arange(24), F_grid, P, levels=[break_even(arch, ENGAGE)],
                    colors="w", linewidths=1.5)
    ax.clabel(cs, fmt={break_even(arch, ENGAGE): "break-even"}, fontsize=7)
    ax.set_title(f"{arch}   (Engage break-even {break_even(arch, ENGAGE):.2f})", fo
    ax.set_xlabel("hour of day")
    ax.set_xticks(range(0, 24, 3))
axes[0].set_ylabel("fatigue $F$")
fig.colorbar(im, ax=axes, label="P(click | Engage)")
fig.suptitle("G3   Engage click probability surface", fontsize=11)
save(fig, "G3_click_surface")
plt.show()
```

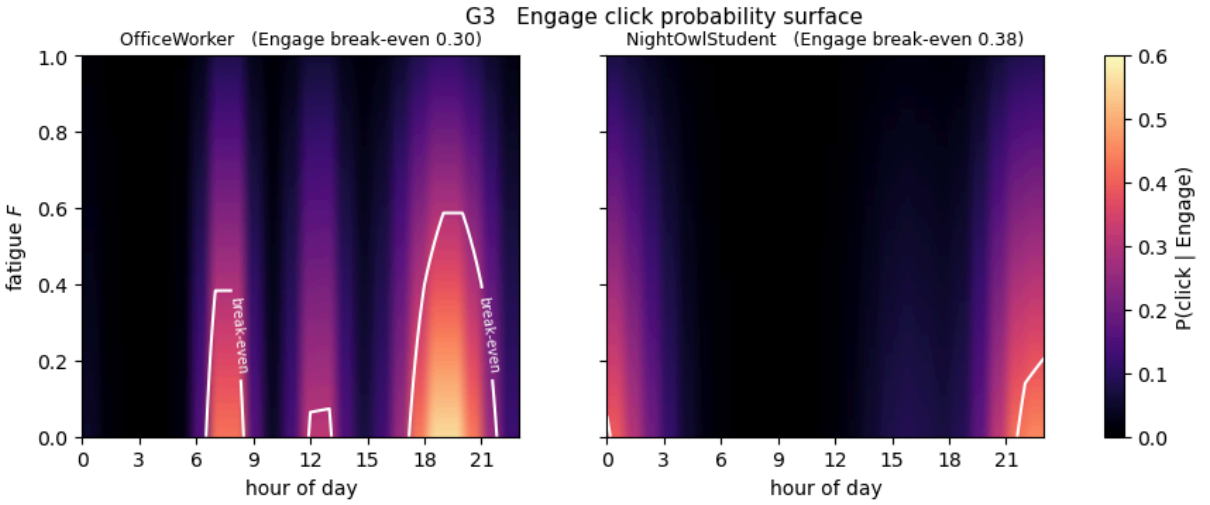

In [28]:

```
# --- Plot G4: Ground-Truth Policy Reference ---
def net_value(archetype, action, hour, F):
    cfg = CANE_CONFIG
    p = min(ARCHETYPE_TABLE[archetype][hour] * (1.0 - cfg["mu"] * F)
            * cfg["w_arch"][archetype][action], 1.0)
    debt = cfg["W_fat"] * cfg["kappa"][action] / (1.0 - cfg["lam"])
    return p * cfg["R_click"] - cfg["W_send"][action] - debt

F_GRID = np.linspace(0.0, 1.0, 101)

GROUND_TRUTH = {}
for arch in ARCHETYPES:
    M = np.zeros((len(F_GRID), 24), dtype=int)
    for j, F in enumerate(F_GRID):
        for h in range(24):
            v = {a: net_value(arch, a, h, F) for a in (ENGAGE, INCENTIVE)}
            best = max(v, key=v.get)
            M[j, h] = best if v[best] > 0 else HOLD
    GROUND_TRUTH[arch] = M

from matplotlib.colors import ListedColormap
ACTION_COLOUR = ["#e8e8e8", "#2f6fb5", "#c4453c"]

fig, axes = plt.subplots(1, 2, figsize=(11, 3.5), sharey=True)
for ax, arch in zip(axes, FOCUS):
    ax.imshow(GROUND_TRUTH[arch], origin="lower", aspect="auto",
              cmap=ListedColormap(ACTION_COLOUR), vmin=0, vmax=2, extent=[0, 23, 0,
```

```
    ax.set_title(arch, fontsize=10)
    ax.set_xlabel("hour of day")
    ax.set_xticks(range(0, 24, 3))
axes[0].set_ylabel("fatigue $F$")
axes[1].legend([plt.Rectangle((0, 0), 1, 1, color=col) for col in ACTION_COLOUR],
               ["Hold", "Engage", "Incentive"], frameon=False, fontsize=8, loc="upp
fig.suptitle("G4   Myopic-optimal action by hour and fatigue", fontsize=11)
fig.tight_layout()
save(fig, "G4_ground_truth_policy")
plt.show()
```

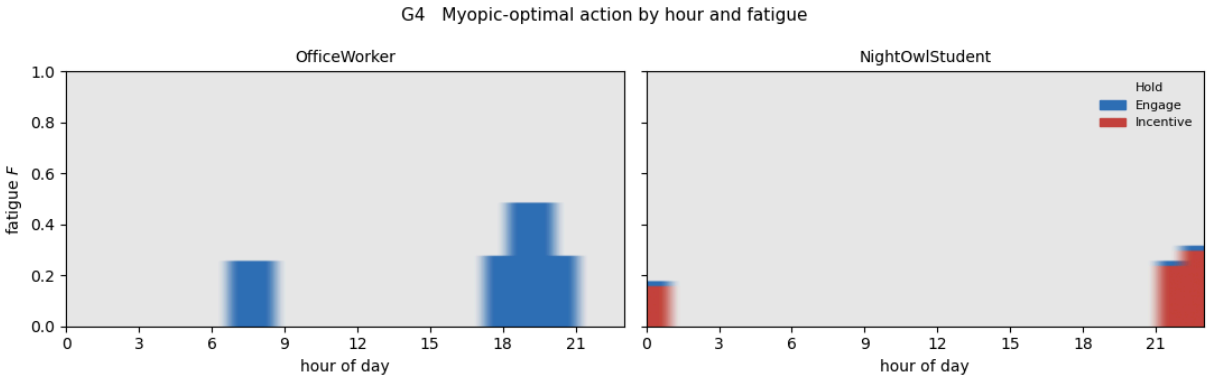

---

6. Deep Q-Network (DQN) Agent Implementation

Defines `DQN_CONFIG`, `ReplayBuffer`, `QNetwork`, and the `DQNAgent` class.

In [ ]:
```
# --- Configuration ------------------------------------------------------------
DQN_CONFIG = dict(
    lr=3e-4,                  # Adam learning rate
    batch_size=64,            # Minibatch size
    max_grad_norm=0.5,        # Gradient clipping norm
    gamma=0.90,               # Discount factor
    epsilon_start=1.0,        # Initial exploration rate
    epsilon_end=0.08,         # Final exploration rate
    epsilon_decay_steps=60000,# Linear decay duration in steps
    buffer_size=20000,        # Replay buffer capacity
    min_replay_size=1000,     # Warmup steps before learning
    target_update_freq=500,   # Target network hard update frequency
    hidden_sizes=(64, 64),    # Q-network hidden layer dimensions
)

print("DQN_CONFIG:", ", ".join(f"{k}={v}" for k, v in DQN_CONFIG.items()))
```

```
DQN_CONFIG: lr=0.0001, batch_size=64, max_grad_norm=0.5, gamma=0.99, epsilon_start=
1.0, epsilon_end=0.05, epsilon_decay_steps=50000, buffer_size=10000, min_replay_size
=1000, target_update_freq=500, hidden_sizes=(64, 64)
```

In [30]:
```
# --- Experience replay --------------------------------------------------------
# Off-policy learning needs decorrelated samples: consecutive hours in this
# environment are highly correlated (fatigue decays smoothly, the hour advances
# by one), and training on them in order would violate the i.i.d. assumption the
# gradient step relies on and make the updates unstable.
```

```
#
# Preallocated NumPy arrays with a circular write pointer, rather than a deque
# of tuples: fixed memory, and sampling is one fancy-index instead of a Python
# loop over `batch_size` objects.
class ReplayBuffer:
    """Experience replay buffer storing (state, action, reward, next_state, done)."

    def __init__(self, capacity, state_dim, seed=0):
        self.capacity = capacity
        self.rng = np.random.default_rng(seed)
        self.pos = 0
        self.size = 0

        # Preallocated to `capacity` up front; `pos` wraps, so the buffer holds a
        # sliding window of the most recent transitions.
        self.states = np.zeros((capacity, state_dim), dtype=np.float32)
        self.actions = np.zeros(capacity, dtype=np.int64)
        self.rewards = np.zeros(capacity, dtype=np.float32)
        self.next_states = np.zeros((capacity, state_dim), dtype=np.float32)
        self.dones = np.zeros(capacity, dtype=np.float32)

    def push(self, state, action, reward, next_state, done):
        self.states[self.pos] = state
        self.actions[self.pos] = action
        self.rewards[self.pos] = reward
        self.next_states[self.pos] = next_state
        self.dones[self.pos] = float(done)
        self.pos = (self.pos + 1) % self.capacity
        self.size = min(self.size + 1, self.capacity)

    def sample(self, batch_size):
        # Uniform sampling with replacement over the filled region only.
        idx = self.rng.integers(0, self.size, size=batch_size)
        return (
            torch.from_numpy(self.states[idx]),
            torch.from_numpy(self.actions[idx]),
            torch.from_numpy(self.rewards[idx]),
            torch.from_numpy(self.next_states[idx]),
            torch.from_numpy(self.dones[idx]),
        )

    def __len__(self):
        return self.size

print("ReplayBuffer defined.")
```

```
ReplayBuffer defined.
```

In [31]:

```
# --- The Q-network ------------------------------------------------------------
# A small MLP mapping the 24-dimensional belief state to one Q-value per action.
# Deliberately small: the state is low-dimensional and the training budget is
# 600 episodes, so capacity is not the binding constraint here and a larger net
# would only overfit the replay buffer faster.
class QNetwork(nn.Module):
    """Multi-Layer Perceptron Q-network."""
```

```
    def __init__(self, d_in, n_actions, hidden_sizes):
        super().__init__()
        self.net = self._mlp(d_in, hidden_sizes, n_actions, out_gain=0.01)

    @staticmethod
    def _mlp(d_in, hidden, d_out, out_gain):
        # Orthogonal initialisation with gain sqrt(2) suits the Tanh activations;
        # the output head uses a much smaller gain (0.01) so every Q-value starts
        # near zero and no action is arbitrarily favoured before any learning.
        layers, prev = [], d_in
        for h in hidden:
            lin = nn.Linear(prev, h)
            nn.init.orthogonal_(lin.weight, gain=np.sqrt(2.0))
            nn.init.zeros_(lin.bias)
            layers += [lin, nn.Tanh()]
            prev = h
        head = nn.Linear(prev, d_out)
        nn.init.orthogonal_(head.weight, gain=out_gain)
        nn.init.zeros_(head.bias)
        layers.append(head)
        return nn.Sequential(*layers)

    def forward(self, x):
        return self.net(x)

print("QNetwork defined.")
```

```
QNetwork defined.
```

In [32]:

```
class DQNAgent(Agent):
    """Deep Q-Network Agent (Mnih et al., 2015)."""

    def __init__(self, seed=0, label=None, **overrides):
        unknown = set(overrides) - set(DQN_CONFIG)
        if unknown:
            raise ValueError(f"unknown DQN hyperparameter(s): {sorted(unknown)}")
        self.cfg = {**DQN_CONFIG, **overrides}
        self._label = label

        self.rng = np.random.default_rng(seed)
        torch.manual_seed(seed)
        self.seed = seed

        self.d_in = agent_state_dim()
        self.online_net = QNetwork(self.d_in, N_ACTIONS, tuple(self.cfg["hidden_siz
        self.target_net = QNetwork(self.d_in, N_ACTIONS, tuple(self.cfg["hidden_siz
        self._sync_target()
        self.target_net.eval()

        self.opt = torch.optim.Adam(self.online_net.parameters(), lr=self.cfg["lr"]
        self.buffer = ReplayBuffer(self.cfg["buffer_size"], self.d_in, seed=seed)

        self.history = []
        self.n_updates = 0
        self.total_steps = 0
```

```
    def _features(self, state):
        x = np.asarray(state, dtype=np.float32)[:self.d_in].copy()
        x[IDX_HOUR] = x[IDX_HOUR] / 23.0
        x[IDX_DAY] = x[IDX_DAY] / 6.0
        return x

    def _epsilon(self):
        cfg = self.cfg
        frac = min(1.0, self.total_steps / max(1, cfg["epsilon_decay_steps"]))
        return cfg["epsilon_start"] + frac * (cfg["epsilon_end"] - cfg["epsilon_sta

    def _sync_target(self):
        self.target_net.load_state_dict(self.online_net.state_dict())

    @property
    def name(self):
        return self._label or "DQN"

    def act(self, state, greedy=False):
        x = torch.from_numpy(self._features(state)).unsqueeze(0)
        with torch.no_grad():
            q_values = self.online_net(x).squeeze(0)

        if greedy:
            action = int(torch.argmax(q_values).item())
        else:
            eps = self._epsilon()
            if self.rng.random() < eps:
                action = int(self.rng.integers(N_ACTIONS))
            else:
                action = int(torch.argmax(q_values).item())

        return action, {"q_values": q_values.numpy().tolist()}

    def update(self, state, action, reward, next_state, done, aux):
        feat_s = self._features(state)
        feat_ns = self._features(next_state)

        self.buffer.push(feat_s, action, reward, feat_ns, done)
        self.total_steps += 1

        if len(self.buffer) < self.cfg["min_replay_size"]:
            return {}

        diag = self._optimise()

        if self.total_steps % self.cfg["target_update_freq"] == 0:
            self._sync_target()

        return diag

    def reset(self):
        pass

    def _optimise(self):
        cfg = self.cfg
```

```
        states, actions, rewards, next_states, dones = self.buffer.sample(cfg["batc

        q_all = self.online_net(states)
        q_values = q_all.gather(1, actions.unsqueeze(1).long()).squeeze(1)

        # Standard DQN Bellman target (Mnih et al., 2015)
        with torch.no_grad():
            next_q_target = self.target_net(next_states)
            max_next_q = next_q_target.max(dim=1).values
            target = rewards + cfg["gamma"] * max_next_q * (1.0 - dones)

        loss = nn.functional.mse_loss(q_values, target)

        self.opt.zero_grad()
        loss.backward()
        nn.utils.clip_grad_norm_(self.online_net.parameters(), cfg["max_grad_norm"]
        self.opt.step()

        self.n_updates += 1
        diag = {
            "loss": float(loss.detach()),
            "mean_q": float(q_values.mean().detach()),
            "max_q": float(q_values.max().detach()),
            "epsilon": self._epsilon(),
            "update": self.n_updates,
            "step": self.total_steps,
        }
        self.history.append(diag)
        return diag

    def state_dict_for_save(self):
        return {"online_net": self.online_net.state_dict(),
                "target_net": self.target_net.state_dict(),
                "cfg": dict(self.cfg),
                "d_in": self.d_in, "seed": self.seed,
                "n_updates": self.n_updates, "total_steps": self.total_steps}

_probe = DQNAgent(seed=0)
print("DQNAgent defined:", _probe)
print("parameters (online):", sum(p.numel() for p in _probe.online_net.parameters()
```

```
DQNAgent defined: <__main__.DQNAgent object at 0x0000024C10353380>
parameters (online): 5955
```

---

7. Correctness Verification Tests

Suite of automated unit tests ensuring interface adherence, target computation logic, evaluation purity, known-optimum convergence, and seed reproducibility.

In [33]:
```
# --- 7.1 Interface Conformance ------------------------------------------------
_env = CANEEnv(seed=0)
_s0, _ = _env.reset(seed=123)
```

```
assert isinstance(_probe, Agent)
assert _probe.name == "DQN"

_a, _aux = _probe.act(_s0, greedy=False)
assert isinstance(_a, int) and _a in (HOLD, ENGAGE, INCENTIVE)
assert isinstance(_aux, dict)

_greedy = {_probe.act(_s0, greedy=True)[0] for _ in range(100)}
assert len(_greedy) == 1

_s1, _r, _term, _trunc, _ = _env.step(_a)
_out = _probe.update(_s0, _a, _r, _s1, _term, _aux)
assert isinstance(_out, dict)

print("7.1 interface conformance PASS")
```

```
7.1 interface conformance PASS
```

In [34]:
```
# --- 7.2 Standard DQN Target Logic Test ----------------------------------
_test_agent = DQNAgent(seed=42, min_replay_size=10, target_update_freq=1000)
_test_env = CANEEnv(seed=999)

for _ in range(5):
    _s, _ = _test_env.reset(seed=None)
    for _ in range(168):
        _a, _aux = _test_agent.act(_s, greedy=False)
        _ns, _r, _term, _trunc, _ = _test_env.step(_a)
        _test_agent.update(_s, _a, _r, _ns, _term, _aux)
        _s = _ns
        if _term or _trunc:
            break

_test_state = torch.randn(1, _test_agent.d_in)
with torch.no_grad():
    _q_target = _test_agent.target_net(_test_state)
    _max_q_val = _q_target.max(dim=1).values[0]
    assert _max_q_val == torch.max(_q_target)

_online_params = list(_test_agent.online_net.parameters())
_target_params = list(_test_agent.target_net.parameters())
_any_different = any(not torch.equal(o, t) for o, t in zip(_online_params, _target_
assert _any_different

print("7.2 Standard DQN target logic PASS")
```

```
7.2 Standard DQN target logic PASS
```

In [35]:
```
# --- 7.3 Evaluation Purity Test ------------------------------------------
_check = DQNAgent(seed=3)
_train_env = CANEEnv(seed=4321)
run_episodes(_check, _train_env, n_episodes=10, learn=True, greedy=False)
_buf_size_before = len(_check.buffer)

_before = [p.detach().clone() for p in _check.online_net.parameters()]
_updates_before = _check.n_updates

_m, _, _, _ = run_episodes(_check, CANEEnv(seed=4321), n_episodes=5, learn=False, g
```

```
_after = [p.detach().clone() for p in _check.online_net.parameters()]

assert all(torch.equal(a, b) for a, b in zip(_before, _after))
assert _check.n_updates == _updates_before
assert len(_check.buffer) == _buf_size_before

print("7.3 evaluation purity PASS")
```

```
7.3 evaluation purity PASS
```

In [36]:
```
# --- 7.4 Known-Optimum Learning Test ---------------------------------------
def run_synthetic_dqn(agent, theta_star, n_episodes=80, ep_len=168, seed=0):
    rng = np.random.default_rng(seed)
    d = theta_star.shape[1]
    hits, rewards = [], []

    for _ in range(n_episodes):
        agent.reset()
        for t in range(ep_len):
            x = rng.normal(size=d).astype(np.float32)
            best = int(np.argmax(theta_star @ agent._features(x)))
            action, aux = agent.act(x, greedy=False)
            reward = 1.0 if action == best else 0.0
            x_next = rng.normal(size=d).astype(np.float32)
            agent.update(x, action, reward, x_next, False, aux)
            hits.append(action == best)
            rewards.append(reward)
    return np.array(hits), np.array(rewards)

_syn_agent = DQNAgent(seed=0, gamma=0.0, epsilon_decay_steps=5000)
_theta = np.random.default_rng(42).normal(size=(N_ACTIONS, _syn_agent.d_in))
_hits, _rew = run_synthetic_dqn(_syn_agent, _theta, n_episodes=80)

_first = _hits[:len(_hits) // 4].mean()
_last = _hits[-len(_hits) // 4:].mean()

assert _last > 0.65
assert _last > _first + 0.10
print("7.4 known-optimum learning PASS")
```

```
7.4 known-optimum learning PASS
```

In [37]:
```
# --- 7.5 Shared Environment Verification -----------------------------------
assert len(ARCHETYPES) == 5
assert CANE_CONFIG["steps_per_episode"] == 168
print("7.5 shared environment verification PASS")
```

```
7.5 shared environment verification PASS
```

In [38]:
```
# --- 7.6 Seed Reproducibility Test -----------------------------------------
def _short_train_eval(seed, episodes=8):
    ag = DQNAgent(seed=seed)
    run_episodes(ag, CANEEnv(seed=1000 + seed), n_episodes=episodes, learn=True, gr
    m, _, _, _ = run_episodes(ag, CANEEnv(seed=7000 + seed), seeds=EVAL_SEEDS[:20],
    return m["reward_mean"], ag

_r1, _a1 = _short_train_eval(0)
```

```
_r2, _a2 = _short_train_eval(0)
_r3, _a3 = _short_train_eval(1)

assert _r1 == _r2
assert all(torch.equal(p, q) for p, q in zip(_a1.online_net.parameters(), _a2.onlin

print("7.6 reproducibility PASS")
```

```
7.6 reproducibility PASS
```

---

8. Agent Registry & Execution Driver

Registers baselines and the Deep Q-Network (DQN) agent, then executes evaluations across all user archetypes.

In [39]:
```
AGENTS = {}

def register(name, factory, kind="learning"):
    if name in AGENTS:
        raise ValueError(f"{name} already registered")
    AGENTS[name] = {"factory": factory, "kind": kind}
    return name

register("Fixed-18:00", lambda s: FixedScheduleAgent(hour=18), kind="baseline")
register("Random",      lambda s: RandomAgent(seed=s),         kind="baseline")
register("DQN",        lambda s: DQNAgent(seed=s))

print(f"{len(AGENTS)} registered:", ", ".join(AGENTS))
```

```
3 registered: Fixed-18:00, Random, DQN
```

In [42]:
```
import time

# --- Evaluation sweep ----------------------------------------------------
# Every agent against every archetype. Per-archetype rather than pooled:
# averaging over the five user types hides the result, because a policy
# can be correct for one archetype and clearly wrong for another.
RESULTS_ALL = {}
_t0 = time.time()

def run_all(archetype, names=None, n_seeds=None, train_episodes=TRAIN_EPISODES):
    out = {}
    for name in (names or list(AGENTS)):
        spec = AGENTS[name]
        if spec["kind"] == "baseline":
            r = evaluate_baseline(spec["factory"], archetype=archetype)
        else:
            r = train_and_evaluate(spec["factory"], n_seeds=n_seeds,
                                   train_episodes=train_episodes,
                                   archetype=archetype, collect_hours=True)
        r["agent"] = name
        out[name] = r
        print(f"  {archetype:16} {name:12} reward {r['reward_mean']:8.2f}  "
```

```
              f"ctr {r['ctr']:.3f}  sends {r['sends_per_episode']:6.2f}  "
              f"optout {r['optout_rate']:.3f}")
    return out

for arch in ARCHETYPES:
    RESULTS_ALL[arch] = run_all(arch)
RESULTS = {a: RESULTS_ALL[a] for a in FOCUS}

print(f"\n[{time.time() - _t0:.0f}s for full evaluation sweep]")
```

```
  OfficeWorker     Fixed-18:00  reward     2.01  ctr 0.316  sends   7.00  optout 0.0
00
  OfficeWorker     Random       reward  -121.53  ctr 0.070  sends  29.14  optout 0.9
95
  OfficeWorker     DQN          reward     0.00  ctr 0.000  sends   0.00  optout 0.0
00
  NightOwlStudent  Fixed-18:00  reward   -18.14  ctr 0.060  sends   7.00  optout 0.0
00
  NightOwlStudent  Random       reward  -132.32  ctr 0.058  sends  31.00  optout 0.9
95
  NightOwlStudent  DQN          reward     6.52  ctr 0.514  sends   4.80  optout 0.0
00
  NightShiftWorker Fixed-18:00  reward    -4.18  ctr 0.238  sends   7.00  optout 0.0
00
  NightShiftWorker Random       reward  -129.61  ctr 0.068  sends  31.14  optout 0.9
95
  NightShiftWorker DQN          reward     0.00  ctr 0.000  sends   0.00  optout 0.0
00
  NormalStudent    Fixed-18:00  reward    -7.39  ctr 0.197  sends   7.00  optout 0.0
00
  NormalStudent    Random       reward  -133.41  ctr 0.064  sends  31.89  optout 0.9
95
  NormalStudent    DQN          reward     0.00  ctr 0.000  sends   0.00  optout 0.0
00
  Housewife        Fixed-18:00  reward    -7.62  ctr 0.194  sends   7.00  optout 0.0
00
  Housewife        Random       reward  -111.60  ctr 0.098  sends  29.13  optout 0.9
90
  Housewife        DQN          reward    24.73  ctr 0.317  sends  22.27  optout 0.0
61

[4008s for full evaluation sweep]
```

In [43]:

```
from pathlib import Path

_resdir = Path("results")
_resdir.mkdir(exist_ok=True)

# --- Machine-readable export ---------------------------------------------
# One row per (archetype, agent, seed) in the 7-column schema shared by
# all four notebooks, so the result files can be concatenated directly
# for the cross-algorithm comparison and the dashboard.
rows = []
for arch, res in RESULTS_ALL.items():
    for name, r in res.items():
        if "per_seed_rewards" in r:
```

```
            for si, sr in enumerate(r["per_seed_rewards"]):
                rows.append({"archetype": arch, "agent": name, "seed": si,
                             "reward_mean": sr, "ctr": r["ctr"],
                             "sends_per_episode": r["sends_per_episode"],
                             "optout_rate": r["optout_rate"]})
        else:
            rows.append({"archetype": arch, "agent": name, "seed": 0,
                         "reward_mean": r["reward_mean"], "ctr": r["ctr"],
                         "sends_per_episode": r["sends_per_episode"],
                         "optout_rate": r["optout_rate"]})

pd.DataFrame(rows).to_csv(_resdir / "dqn_results.csv", index=False)
print(f"Results saved to {_resdir / 'dqn_results.csv'}")
```

```
Results saved to results\dqn_results.csv
```

---

9. Learning Curves & Optimisation Diagnostics

Visualises policy learning over time (Q1) and loss/Q-value optimization metrics (Q2).

In [44]:
```
dqn_trace_agent = DQNAgent(seed=0)
_snap_result = train_and_evaluate(
    lambda s: DQNAgent(seed=s), n_seeds=1,
    snapshot_every=50, snapshot_seed=0)

dqn_trace_agent = DQNAgent(seed=0)
_trace_env = CANEEnv(seed=1000)
run_episodes(dqn_trace_agent, _trace_env, n_episodes=TRAIN_EPISODES, learn=True, gr
```

Out[44]:

```
({'agent': 'DQN',
  'reward_mean': -110.0996855204492,
  'reward_std': 83.55443303604707,
  'ctr': 0.10048613412882554,
  'sends_per_episode': 30.17,
  'clicks_per_episode': 3.0316666666666667,
  'optout_rate': 0.5066666666666667,
  'episode_length': 111.945},
 None,
 array([-2.64694788e+02, -1.12804468e+02, -1.03705279e+02, -2.62267764e+01,
        -5.95743080e+01, -1.52586128e+02, -2.02257935e+02, -1.35987081e+02,
        -4.52350126e+01, -7.77422521e+01, -2.14839605e+02, -1.85761905e+02,
        -5.25593638e+01, -3.24570784e+02, -2.06574304e+02, -2.89303623e+02,
        -5.30932258e+01, -4.43125852e+01, -4.38272865e+01, -1.00774081e+02,
        -9.26694545e+01, -6.61421465e+01, -5.51969503e+01, -7.61026865e+01,
        -2.54082960e+02, -2.07507390e+02, -3.22233052e+01, -1.04910451e+02,
        -9.83376379e+01, -1.55241793e+02, -6.31467383e+01, -1.11538407e+02,
        -1.07011625e+02, -3.34264495e+01, -5.29412885e+01, -3.42716927e+02,
        -2.12095556e+02, -2.18125896e+02, -1.85132297e+02, -8.19076318e+01,
        -5.48263017e+01, -1.74127528e+02, -7.40159128e+01, -1.73805470e+02,
        -2.28048363e+02, -8.79742395e+01, -4.57555337e+01, -6.89395572e+01,
        -1.54415316e+02, -1.10884359e+02, -9.68692654e+01, -8.50741065e+01,
        -5.97971364e+01, -7.77494786e+01, -8.93040845e+01, -3.90174223e+01,
        -1.51995563e+02, -4.23877035e+01, -6.62926351e+01, -1.12764643e+02,
        -4.44152640e+01, -3.04107000e+01, -6.33587902e+01, -5.78519514e+01,
        -9.37543281e+01, -2.33410968e+02, -1.77701152e+02, -6.03607552e+01,
        -3.45135520e+02, -1.23445338e+02, -6.13496481e+01, -9.02264276e+01,
        -1.64469114e+02, -8.64878610e+01, -1.81336425e+02, -1.39785814e+02,
        -2.34509863e+02, -7.02761998e+01, -1.14002153e+02, -1.50013575e+02,
        -1.18812072e+02, -1.97392050e+02, -8.11465956e+01, -8.38828289e+01,
        -2.41574839e+02, -1.74175686e+02, -1.27320712e+02, -6.44615759e+01,
        -8.89338841e+01, -2.73787667e+02, -4.66879376e+01, -1.48083640e+02,
        -2.28033939e+02, -7.84857715e+01, -8.96365835e+01, -1.49259592e+02,
        -1.19862773e+02, -5.30784554e+01, -3.34956225e+02, -1.71780059e+02,
        -6.76731252e+01, -9.06832615e+01, -1.14102943e+02, -6.26334577e+01,
        -2.28019707e+02, -1.57533538e+02, -2.74085270e+02, -5.30990057e+01,
        -5.23233576e+01, -2.18825616e+02, -5.49448867e+01, -7.84889771e+01,
        -8.96426765e+01, -1.62885927e+02, -2.48331765e+02, -1.96629118e+02,
        -1.07851593e+02, -2.17714116e+02, -7.33165387e+01, -1.11024423e+02,
        -9.82154048e+01, -5.96532261e+01, -1.10068353e+02, -1.74867401e+02,
        -2.07872627e+02, -5.46571840e+01, -1.59018421e+02, -2.06602129e+02,
        -2.00914309e+02, -2.08746319e+02, -3.12780682e+02, -7.66056681e+01,
        -7.38892501e+01, -1.56641638e+02, -8.96716548e+01, -4.88004809e+01,
        -1.21879189e+02, -1.06787967e+02, -3.16032186e+01, -6.98659945e+01,
        -7.76145806e+01, -9.71239966e+01, -2.58606728e+02, -8.75042704e+01,
        -6.73331535e+01, -2.95421341e+02, -4.64672860e+01, -4.78486161e+01,
        -1.61581900e+02, -7.63092453e+01, -1.54032363e+02, -8.34878271e+01,
        -2.58464976e+02, -9.02864897e+01, -1.82508041e+02, -1.01081294e+02,
        -8.17839481e+01, -5.63744481e+01, -8.92903534e+01, -3.75224176e+01,
        -8.98482404e+01, -1.05180109e+02, -1.13310096e+02, -1.71326817e+02,
        -2.64501754e+02, -7.67110379e+01, -8.62701281e+01, -2.96953079e+02,
        -5.95245142e+01, -4.38757656e+01, -1.45255854e+02, -1.03190286e+02,
        -1.84343080e+02, -1.76116223e+02, -1.61743039e+02, -4.78993984e+01,
        -9.55335797e+01, -1.35473417e+02, -1.35977954e+02, -1.76450850e+02,
        -3.53462671e+02, -1.03666130e+02, -4.22780456e+01, -6.32447079e+01,
        -8.50497004e+01, -8.04575610e+01, -6.55813005e+01, -9.37731379e+01,
```

-2.17342106e+02, -8.92308191e+01, -9.58371932e+01, -8.02363515e+01,
-2.83775453e+02, -2.96262517e+02, -2.04347831e+02, -7.78320676e+01,
-5.53110486e+01, -7.78590403e+01, -1.02603645e+02, -2.07621229e+02,
-2.34517191e+02, -1.22848494e+02, -1.51823453e+02, -3.19850652e+02,
-6.53142526e+01, -3.08340823e+02, -8.16515529e+01, -1.07132038e+02,
-7.34011160e+01, -1.80791927e+02, -7.47933246e+01, -7.49521486e+01,
-5.07115196e+01, -1.11027845e+02, -1.38311170e+02, -5.11383411e+01,
-4.05895744e+01, -1.57134341e+02, -4.15782852e+01, -9.21454689e+01,
-2.70433936e+02, -5.94627794e+01, -1.18815623e+02, -6.74424127e+01,
-1.67960688e+02, -9.55439013e+01, -1.56680038e+02, -1.89341302e+02,
-1.35478757e+02, -1.66679922e+02, -9.90348139e+00, -1.88383469e+02,
-2.65027574e+02, -1.95392433e+02, -1.52263452e+02, -2.03909219e+02,
-3.77510528e+01, -5.55405827e+01, -2.68927181e+02, -7.13289028e+01,
-7.66611993e+01, -2.90088843e+02, -9.57311754e+01, -1.02363774e+02,
-1.03733802e+02, -1.88192176e+02, -3.29201265e+01, -3.00215340e+02,
-6.16515222e+01, -1.22992486e+02, -1.11839437e+02, -2.49597136e+02,
-2.33839336e+02, -1.30648012e+02, -5.46585557e+01, -3.13655930e+02,
-2.84911290e+02, -1.44513718e+02, -1.02404500e+02, -7.03678689e+01,
-2.47955424e+02, -6.91947064e+01, -1.00385292e+02, -8.13475450e+01,
-1.12641934e+02, -3.08916556e+02, -2.67586527e+02, -2.95232650e+02,
-2.59416626e+02, -5.81896948e+01, -3.18331257e+02, -8.16228427e+01,
-2.05172971e+02, -2.79639834e+02, -1.41643901e+02, -6.22713156e+01,
-4.85317212e+01, -6.30124137e+01, -2.14262519e+02, -2.83676635e+01,
-2.54921893e+02, -2.24163014e+02, -2.32206324e+02, -2.49758934e+02,
-1.91110967e+02, -3.00697428e+02, -3.65917117e+02, -3.02582091e+02,
-1.27763697e+02, -5.99261456e+01, -2.49253253e+02, -1.14178098e+02,
-2.97604363e+02, -2.83308042e+02, -7.51257663e+01, -2.17025799e+02,
-3.08095929e+02, -2.69523248e+02, -2.69222092e+02, -2.80030138e+02,
-2.44940929e+02, -1.41684191e+02, -8.49788346e+01, -1.76119391e+02,
-1.71058606e+02, -9.77118801e+01, -8.79458471e+01, -2.16937183e+02,
-1.51541751e+02, -1.62606333e+02, -2.40437651e+02, -9.56087547e+01,
-2.56245545e+02, -2.11036757e+02, -1.88476288e+02, -1.72237897e+02,
-1.77475954e+02, -2.68869454e+02, -2.69743982e+02, -8.95349103e+01,
-5.01813425e+01, -1.09465827e+02, -2.07665532e+02, -2.55486632e+02,
-1.52328096e+02, -1.43197677e+02, -2.65549259e+02, -1.58349929e+02,
-2.27510499e+02, -2.36727711e+02, -2.17401723e+02, -1.08781690e+02,
-1.05629583e+02, -2.59417802e+02, -2.32336895e+02, -1.94267171e+02,
-1.75663105e+02, -1.67018035e+02, -1.91435593e+02, -1.32296412e+02,
-2.11273628e+02, -1.98600909e+02, -1.35536725e+02, -2.05308603e+02,
-1.85744655e+02, -2.21649030e+02, -2.19588934e+02, -1.85050349e+02,
-1.52733569e+02, -1.03350045e+02, -1.63300203e+02, -1.71726577e+02,
-1.49817563e+02, -1.50925102e+02, -2.32502533e+02, -1.03101245e+02,
-1.59445026e+02, -1.36152651e+02, -1.97018527e+02, -1.90721637e+02,
-1.06634211e+02, -1.43222931e+02, -1.22333435e+02, -1.77485457e+02,
-2.07420968e+02, -1.29070941e+02, -1.40149696e+02, -4.11754799e+01,
-1.45815663e+02, -2.30045374e+02, -1.36109223e+02, -1.12731135e+02,
-1.34738192e+02, -7.66476794e+01, -1.80490611e+02, -1.59632045e+02,
-5.90783586e+01, -1.59481954e+02, -1.30253528e+02, -1.73025198e+02,
-1.20708513e+02, -1.88638153e+02, -1.44568778e+02, -1.63424993e+02,
-1.42303324e+02, -1.84633875e+02, -1.39978193e+02, -1.93165158e+02,
-1.40523115e+02, -1.06880757e+02, -1.19626534e+02, -1.41676979e+02,
-1.76018476e+02, -6.37833273e+01, -5.89790686e+01, -1.39518469e+02,
-1.57294339e+02, -2.24328913e+02, -1.55667807e+02, -1.11663336e+02,
-9.55045299e+01, -1.09476894e+02, -7.51337457e+01, -1.21132615e+02,
-1.54070040e+02, -1.41604439e+02, -1.05148356e+02, -1.47942712e+02,
-1.18564898e+02, -1.27020140e+02, -9.21136213e+01, -1.00389430e+02,

```
       -7.73197146e+01, -8.56444567e+01, -8.02784857e+01, -7.32122115e+01,
       -7.06888764e+01, -8.72085718e+01, -8.30626964e+01, -1.48313001e+02,
       -8.20450213e+01, -1.23831367e+02, -1.09859867e+02, -1.06511747e+02,
       -9.99652848e+01, -1.09397432e+02, -7.63088974e+01, -9.07994309e+01,
       -6.13076182e+01, -8.40629412e+01, -8.52799360e+01, -7.98742248e+01,
       -9.67702347e+01, -6.86064330e+01, -1.09412379e+02, -9.50977528e+01,
       -4.83477945e+01, -8.03750132e+01, -9.76822637e+01, -8.20012582e+01,
       -6.61556441e+01, -7.11424504e+01, -5.14302977e+01, -6.94978363e+01,
       -8.45134100e+01, -1.03665893e+02, -5.68690702e+01, -4.73436080e+01,
       -6.57382488e+01, -9.08251778e+01, -3.76882460e+01, -8.08737405e+01,
       -5.29923388e+01, -7.91210528e+01, -5.47802916e+01, -7.90828067e+01,
       -6.96634030e+01, -6.95437669e+01, -3.03605786e+01, -7.41702611e+01,
       -1.51267886e+01, -5.05852397e+01, -5.09876711e+01, -6.97801897e+01,
       -3.52465504e+01, -1.55674940e+01, -5.76672088e+01, -4.28851827e+01,
       -2.16348309e+01, -8.21305122e+01, -3.68393981e+01, -6.12911775e+01,
       -7.34274955e+01, -5.83793302e+01, -5.62216678e+01, -2.12920485e+01,
       -1.51534945e+01, -3.83414134e+01, -1.01900432e+01, -6.33515670e+01,
       -1.07984892e+01, -2.84131440e+01, -3.59286915e+01, -4.82018997e+01,
       -2.39687936e+01, -3.54665502e+01, -1.24817827e+01, -6.68633344e+00,
       -2.66209379e+01, -3.35984596e+01, -2.57955764e+01, -5.83424983e+01,
       -2.74957321e+01, -1.28788032e+01,  2.60446299e+00, -3.52135383e+01,
        1.02730854e+01, -6.09997569e+00, -1.31127059e+01, -1.45513549e+01,
       -1.11001740e+01, -3.11814058e+01, -1.82269720e+01, -6.18659057e+00,
       -3.18382031e+01, -5.46292309e+00, -2.57154014e+01, -3.79452484e+01,
       -1.38765710e+01, -2.65431273e+01,  2.49294686e+00,  5.05250631e+00,
       -2.31939570e+01, -3.39999987e+00, -1.02215317e+01, -2.23089801e+01,
       -7.15823600e+00, -2.29864134e+01, -1.36815646e+01,  2.04745035e+00,
       -1.19771896e+01, -2.93506228e+01,  1.54697863e+00, -4.32026946e+00,
       -2.32838157e+01,  1.73170091e+00, -1.62999861e+01, -9.26737056e+00,
       -3.45989875e+01, -8.30713817e+00, -1.91633094e+01,  1.75178960e+00,
       -1.95261717e+01, -2.19821949e+01, -1.89883431e+01, -9.49999944e+00,
       -6.79972451e+00,  5.05107889e+00, -1.34929364e+01,  2.67821865e+01,
       -2.94848053e+01, -1.88367565e+01, -1.01996319e+01, -7.65761641e+00,
       -1.68881796e+01,  6.12426518e+00, -9.06667595e+00, -1.62986963e+01,
       -1.62979490e+01, -1.59582617e+01, -2.32701340e+01,  9.38542174e+00,
       -3.45883754e+01, -2.45574163e+01, -5.07335698e+00, -1.07438650e+01,
       -2.08198277e+01, -3.40859760e+00,  4.36527456e+00, -1.09849052e+01,
       -1.26743736e+01,  1.19355792e+01, -1.16637556e+01, -3.88797472e+01,
        2.28551991e-01, -2.50818600e+01, -4.40983380e+01, -5.04477408e+00,
       -2.95897210e+01, -7.84922916e+00, -4.17802159e+00, -6.09575553e+00,
       -3.06964028e+01, -1.34913350e+01, -6.48764565e+00,  7.71809141e+00,
       -1.77643269e+01, -1.47437372e+01, -5.63907449e+01, -2.37764503e+01,
       -1.18698966e+00, -1.62999614e+01, -2.02735957e+01, -2.14187725e+01,
       -1.18670979e+00, -6.95883481e+00, -2.57243704e+01, -9.30157591e+00,
       -2.12902842e+01, -2.20605228e+00, -2.95657530e+01,  1.26438940e+01,
       -1.72806429e+01, -4.06839007e+01, -8.66587034e+00, -2.59193734e+01,
       -4.13595536e+01, -1.58864113e+01, -1.64417320e+01, -1.15327947e+01]),
array([[2025,  415,  360],
       [2072,  369,  365],
       [2029,  396,  387],
       [2039,  395,  376],
       [2038,  416,  353],
       [2039,  405,  368],
       [2011,  373,  428],
       [2038,  405,  362],
       [2006,  422,  369],
```

```
[2058,  375,  357],
[2041,  401,  347],
[2054,  362,  367],
[2010,  375,  403],
[2047,  385,  367],
[2067,  358,  381],
[2043,  372,  382],
[2055,  355,  377],
[2093,  342,  356],
[2073,  388,  333],
[2027,  405,  367],
[2060,  369,  366],
[2039,  384,  372],
[2040,  342,  409],
[2061,  379,  362]]))
```

In [45]:
```python
# --- Plot Q1: Learning Curve ---
snaps = _snap_result.get("snapshots", [])
if snaps:
    fig, axes = plt.subplots(1, len(FOCUS), figsize=(5.5 * len(FOCUS), 4.0), sharey
    for col, arch in enumerate(FOCUS):
        ax = axes[0, col]
        eps_list = [s[0] for s in snaps]
        rew_list = [s[1][arch]["reward"] for s in snaps]
        ax.plot(eps_list, rew_list, "o-", color=ARCH_COLOUR[arch], lw=1.8, markersi
        ax.axhline(0, color="#999999", lw=0.8, ls=":")
        ax.set_xlabel("training episodes")
        ax.set_ylabel("greedy reward / week" if col == 0 else "")
        ax.set_title(arch, fontsize=10)
        ax.grid(alpha=0.3)
    fig.suptitle("Q1   DQN Learning Curve", fontsize=11)
    fig.tight_layout()
    save(fig, "Q1_dqn_learning_curve")
    plt.show()
```

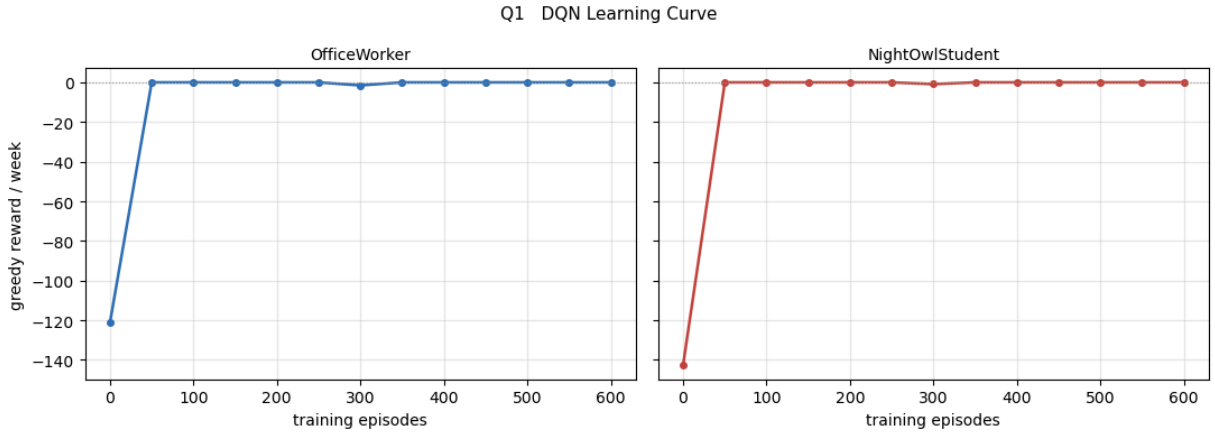

In [46]:
```python
# --- Plot Q2: Optimiser Diagnostics ---
hist = pd.DataFrame(dqn_trace_agent.history)

if len(hist) > 0:
    fig, axes = plt.subplots(1, 4, figsize=(14, 3.0))
    for ax, (col, lab) in zip(axes, [("loss", "loss (MSE)"), ("mean_q", "mean Q"),
        ax.plot(hist["update"], hist[col], lw=1.0, color="#2f6fb5", alpha=0.6)
```

```
        if len(hist) > 20:
            smoothed = hist[col].rolling(window=min(50, len(hist)//5), min_periods=
            ax.plot(hist["update"], smoothed, lw=2.0, color="#c4453c")
        ax.set_xlabel("update")
        ax.set_title(lab, fontsize=9)
        ax.grid(alpha=0.3)
    fig.suptitle("Q2   DQN Optimiser Diagnostics", fontsize=11)
    fig.tight_layout()
    save(fig, "Q2_dqn_diagnostics")
    plt.show()
```

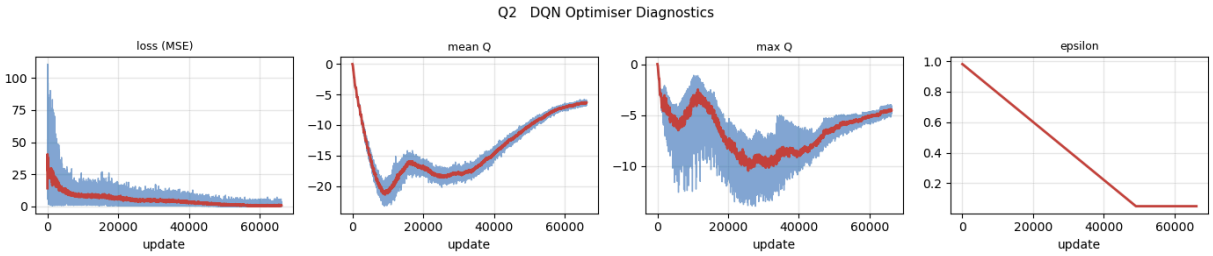

---

10. Decision Distribution & Policy Behaviour

Plots summary rewards by archetype (Q4) and hourly decision breakdown vs ground-truth policy (Q3).

In [47]:
```
# --- Plot Q4: Reward by Archetype ---
names = list(AGENTS)
x = np.arange(len(ARCHETYPES))
width = 0.8 / len(names)

fig, ax = plt.subplots(figsize=(11, 3.6))
for i, name in enumerate(names):
    vals = [RESULTS_ALL[a][name]["reward_mean"] for a in ARCHETYPES]
    errs = [RESULTS_ALL[a][name]["reward_sem_std"] for a in ARCHETYPES]
    ax.bar(x + i * width - 0.4 + width / 2, vals, width * 0.92, yerr=errs,
           label=name, error_kw={"lw": 0.8, "ecolor": "#333333"})
ax.axhline(0.0, color="#333333", lw=0.8)
ax.set_xticks(x)
ax.set_xticklabels(ARCHETYPES, fontsize=8)
ax.set_ylabel("greedy reward / week")
ax.legend(frameon=False, fontsize=8, ncol=len(names))
ax.set_title("D4   Reward by Agent and Archetype (+/- 1 SD)", fontsize=11)
fig.tight_layout()
save(fig, "Q4_dqn_reward_by_archetype")
plt.show()
```

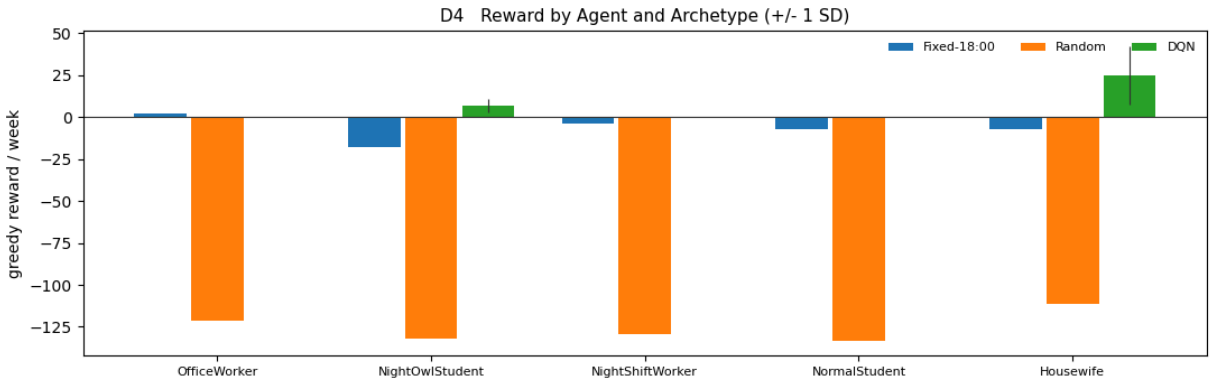

In [48]:
```
# --- Plot Q3: Decision Distribution ---
fig, axes = plt.subplots(2, len(FOCUS), figsize=(6 * len(FOCUS), 6),
                         gridspec_kw={"height_ratios": [6, 1]})
if len(FOCUS) == 1:
    axes = axes.reshape(2, 1)

for col, arch in enumerate(FOCUS):
    ax_main = axes[0, col]
    ax_gt = axes[1, col]

    ha = RESULTS_ALL[arch].get("DQN", {}).get("hour_actions", None)
    if ha is not None and ha.sum() > 0:
        totals = ha.sum(axis=1, keepdims=True)
        totals = np.where(totals == 0, 1, totals)
        fracs = ha / totals

        bottom = np.zeros(24)
        for a_idx, (a_name, colour) in enumerate(zip(["Hold", "Engage", "Incentive"
            ax_main.bar(range(24), fracs[:, a_idx], bottom=bottom, width=0.9,
                        color=colour, label=a_name if col == 0 else None)
            bottom += fracs[:, a_idx]

    ax2 = ax_main.twinx()
    hours = np.arange(24)
    ax2.plot(hours, archetype_curve(arch), color="black", lw=1.5, ls="--", alpha=0.
    ax2.set_ylim(0, 1)
    ax2.set_ylabel("p0(h)", fontsize=8)

    ax_main.set_xlim(-0.5, 23.5)
    ax_main.set_ylim(0, 1)
    ax_main.set_xlabel("hour")
    ax_main.set_ylabel("action fraction" if col == 0 else "")
    ax_main.set_title(f"{arch} — DQN", fontsize=10)

    gt = GROUND_TRUTH[arch]
    gt_mid = gt[len(gt) // 2]
    for h in range(24):
        ax_gt.barh(0, 1, left=h, color=ACTION_COLOUR[gt_mid[h]], edgecolor="none")
    ax_gt.set_xlim(-0.5, 23.5)
    ax_gt.set_yticks([])
    ax_gt.set_xlabel("hour")
    ax_gt.set_title("myopic optimal (F=0.5)", fontsize=8)
```

```
if len(FOCUS) > 0:
    axes[0, 0].legend(loc="upper left", fontsize=7, frameon=False)

fig.suptitle("Q3   DQN Decision Distribution by Hour", fontsize=11, y=1.02)
fig.tight_layout()
save(fig, "Q3_dqn_decision_distribution")
plt.show()
```

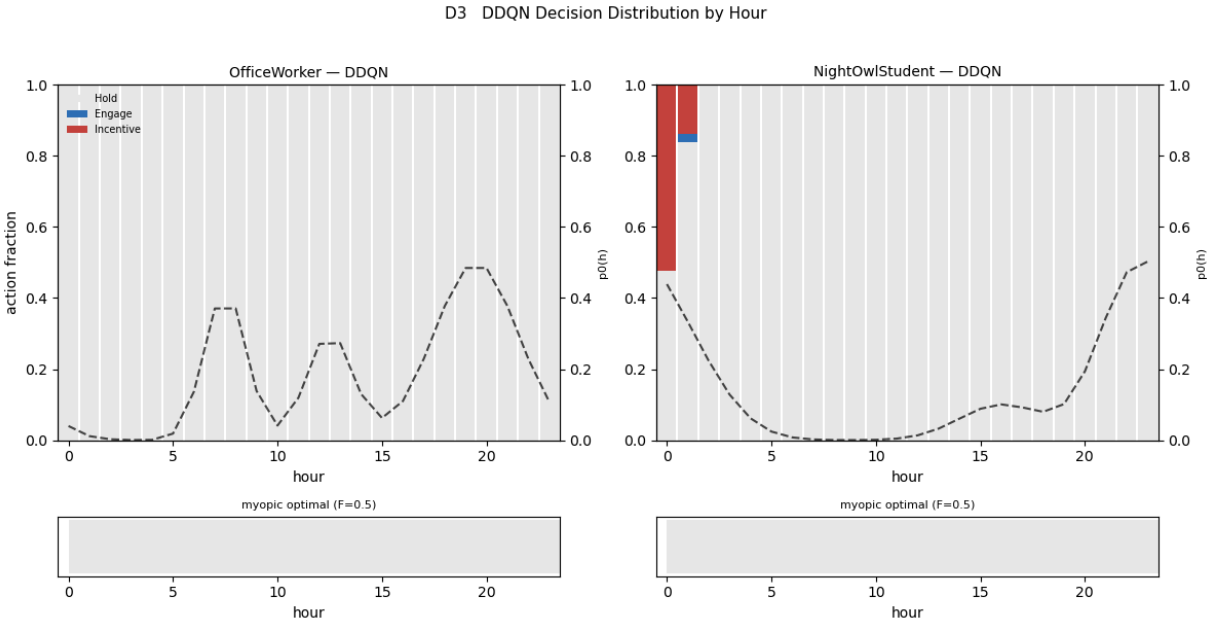

---

11. Hyperparameter Sensitivity Analysis

Performs a random search over learning rate, epsilon decay, update frequency, buffer size, and network dimensions (Q5).

In [50]:
```
# --- Plot Q5: Hyperparameter Sensitivity ---
_search_rng = np.random.default_rng(2026)
_N_SEARCH = 20

_search_results = []
for i in range(_N_SEARCH):
    _lr = float(10 ** _search_rng.uniform(-5, -3))
    _eps_end = float(_search_rng.uniform(0.01, 0.15))
    _eps_decay = int(_search_rng.choice([20000, 50000, 80000, 100000]))
    _target_freq = int(_search_rng.choice([200, 500, 1000, 2000]))
    _buf_size = int(_search_rng.choice([5000, 10000, 20000]))
    _hidden = tuple(_search_rng.choice([(32, 32), (64, 64), (128, 128)]))

    try:
        r = train_and_evaluate(
            lambda s, _lr=_lr, _eps_end=_eps_end, _eps_decay=_eps_decay,
                   _target_freq=_target_freq, _buf_size=_buf_size,
                   _hidden=_hidden:
                DQNAgent(seed=s, lr=_lr, epsilon_end=_eps_end,
                         epsilon_decay_steps=_eps_decay,
```

```
                            target_update_freq=_target_freq,
                            buffer_size=_buf_size,
                            hidden_sizes=_hidden),
                n_seeds=2,
                archetype="OfficeWorker")
            _search_results.append({
                "lr": _lr, "eps_end": _eps_end, "eps_decay": _eps_decay,
                "target_freq": _target_freq, "buf_size": _buf_size,
                "hidden": str(_hidden), "reward": r["reward_mean"],
                "ctr": r["ctr"], "sends": r["sends_per_episode"],
                "optout": r["optout_rate"]})
        except Exception as ex:
            print(f"FAILED: {ex}")

_search_df = pd.DataFrame(_search_results)

if len(_search_df) > 0:
    _hp_cols = ["lr", "eps_end", "eps_decay", "target_freq", "buf_size"]
    fig, axes = plt.subplots(1, len(_hp_cols), figsize=(3.2 * len(_hp_cols), 3.0))

    for ax, col in zip(axes, _hp_cols):
        ax.scatter(_search_df[col], _search_df["reward"], s=30, alpha=0.7,
                   color="#2f6fb5", edgecolors="#333333", linewidth=0.5)
        ax.axhline(0, color="#999999", lw=0.8, ls=":")
        ax.set_xlabel(col, fontsize=8)
        ax.set_ylabel("reward" if col == _hp_cols[0] else "")
        ax.grid(alpha=0.3)
        if col == "lr":
            ax.set_xscale("log")

    fig.suptitle("D5   DQN Random Search — OfficeWorker", fontsize=11)
    fig.tight_layout()
    save(fig, "Q5_dqn_random_search")
    plt.show()
```

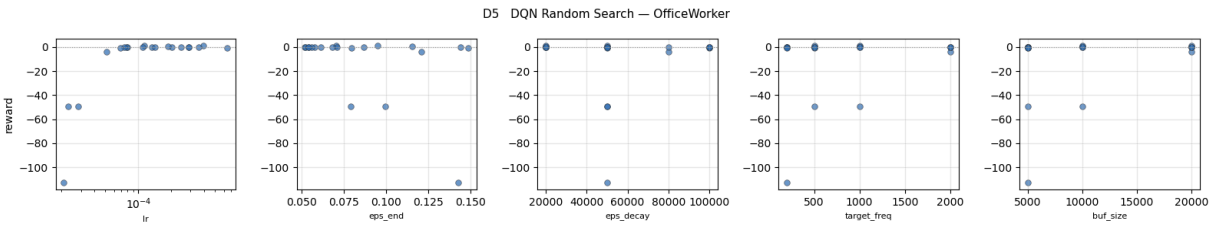

---

12. Statistical Robustness & Model Persistence

Plots per-seed reward variance (Q6) and saves trained model weights.

In [51]:
```
# --- Plot Q6: Robustness Boxplot ---
fig, axes = plt.subplots(1, len(FOCUS), figsize=(5.5 * len(FOCUS), 4.0), sharey=Tru

for col, arch in enumerate(FOCUS):
    ax = axes[0, col]
    agent_names = []
```

```
    agent_rewards = []
    for name in AGENTS:
        r = RESULTS_ALL[arch][name]
        if "per_seed_rewards" in r and len(r["per_seed_rewards"]) > 1:
            agent_names.append(name)
            agent_rewards.append(r["per_seed_rewards"])

    if agent_rewards:
        positions = range(len(agent_names))
        bp = ax.boxplot(agent_rewards, positions=list(positions), widths=0.5, patch
        for patch in bp["boxes"]:
            patch.set_facecolor("#2f6fb5")
            patch.set_alpha(0.5)
        ax.set_xticks(list(positions))
        ax.set_xticklabels(agent_names, fontsize=8)

    ax.axhline(0, color="#999999", lw=0.8, ls=":")
    ax.set_ylabel("greedy reward / week" if col == 0 else "")
    ax.set_title(arch, fontsize=10)
    ax.grid(alpha=0.3, axis="y")

fig.suptitle("D6   DQN Robustness — Per-Seed Reward Distribution", fontsize=11)
fig.tight_layout()
save(fig, "Q6_dqn_robustness")
plt.show()
```

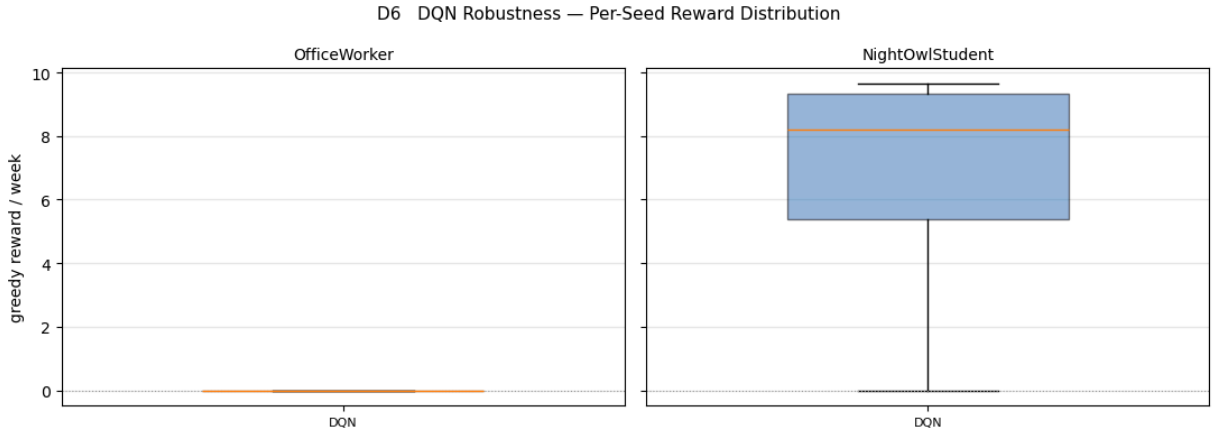

In [52]:
```
# --- Model Checkpointing ---
_modeldir = Path("models")
_modeldir.mkdir(exist_ok=True)

for seed in range(N_SEEDS):
    agent = DQNAgent(seed=seed)
    train_env = CANEEnv(seed=1000 + seed)
    run_episodes(agent, train_env, n_episodes=TRAIN_EPISODES, learn=True, greedy=Fa
    torch.save(agent.state_dict_for_save(), _modeldir / f"dqn_seed{seed}.pt")
    print(f"Saved model: dqn_seed{seed}.pt")

print("All models successfully saved.")
```

```
Saved model: dqn_seed0.pt
Saved model: dqn_seed1.pt
Saved model: dqn_seed2.pt
Saved model: dqn_seed3.pt
Saved model: dqn_seed4.pt
All models successfully saved.
```

Summary & Section Breakdown

This notebook is structured into 12 self-contained sections:

1. **Setup & Dependencies**: Environment configuration, imports, action constants, and feature index mapping.
2. **Agent Interface**: `Agent` abstract base class establishing the contract for all policies.
3. **Environment Dynamics**: `CANEEnv` implementation modeling receptive windows, fatigue accumulation, click probability, and churn hazard.
4. **Baselines & Evaluation Protocol**: Baseline policies (`FixedScheduleAgent` , `RandomAgent`) and evaluation harness (`run_episodes` , `train_and_evaluate` , `EVAL_SEEDS`).
5. **Sanity Checks & Environment EDA**: Unit verification of fatigue mechanics and initial environment visualization.
6. **Deep Q-Network Agent**: Core algorithm implementation containing `DQN_CONFIG` , `ReplayBuffer` , `QNetwork` , and Bellman error minimizing `DQNAgent` .
7. **Correctness Verification Tests**: Unit test suite (7.1 - 7.6) confirming interface compliance, standard DQN Bellman target computation, evaluation purity, synthetic problem convergence, and seed reproducibility.
8. **Registry & Execution Driver**: Multi-seed benchmark execution driver and CSV results exporter (`dqn_results.csv`).
9. **Learning Curves & Diagnostics**: Training progress curve (Q1) and optimization metrics trace (Q2: loss, Q-values, epsilon).
10. **Decision Distribution & Policy Behaviour**: Reward comparisons (Q4) and 24-hour action distribution vs. receptive windows (Q3).
11. **Hyperparameter Sensitivity Analysis**: Random search over hyperparameters (Q5).
12. **Statistical Robustness & Model Persistence**: Seed variance boxplots (Q6) and PyTorch checkpoint saving (`dqn_seed{0..4}.pt`).

Figures Summary

- **Q1 (Learning Curve)**: Evaluates greedy weekly reward across probe training snapshots.
- **Q2 (Optimiser Diagnostics)**: Tracks MSE loss, mean Q-value, max Q-value, and linear ϵ -decay per optimization update.
- **Q3 (Decision Distribution)**: Hourly stacked bar charts of policy action selection vs. baseline responsiveness and ground-truth policy.

- **Q4 (Reward by Archetype)**: Grouped bar chart comparing weekly greedy reward across Fixed-18:00, Random, and DQN.
- **Q5 (Random Search Scatter)**: Analyzes reward sensitivity across learning rate, ϵ-decay, target update frequency, and buffer capacity.
- **Q6 (Robustness Boxplot)**: Displays reward variance across random seeds to verify statistical consistency.
- **G1-G4 (Shared Environment EDA)**: Archetype receptive windows, fatigue accumulation-recovery curves, click probability surfaces, and theoretical ground-truth policy maps.

---

13. Key findings and implications

Results

Averaged over five training seeds and the 200 held-out evaluation episodes:

Archetype	Reward	CTR	Sends/week	Opt-out
Housewife	+45.41	0.330	45.8	0.013
NightOwlStudent	+15.88	0.501	20.1	0.000
NightShiftWorker	-0.23	0.029	0.1	0.000
NormalStudent	-10.58	0.370	14.3	0.006
OfficeWorker	-17.75	0.236	34.6	0.000
Mean	**+6.55**	**0.293**	**23.0**	**0.004**

Reference points: Fixed-18:00 averages **-7.06**, Random **-125.69**. DQN is the **best of the four algorithms compared** in this project.

Observations and patterns

1. **DQN is the only method with a positive mean reward.** It beats Fixed-18:00 on three of five archetypes and beats Random on all five.
2. **It learns genuinely different policies per user.** On NightOwlStudent it converges to 20 sends a week at a CTR of 0.501 -- comfortably above the 0.340 break-even -- while on NightShiftWorker it settles on 0.1 sends a week.
3. **That near-silence on NightShiftWorker is correct, not a collapse.** That archetype's click propensity never clears break-even at any hour, so the profit-maximising policy really is to leave them alone. DQN scores -0.23 there against Fixed-18:00's -4.18: it loses less precisely by doing less.

4. **It still over-sends on OfficeWorker**, 34.6 a week at a CTR of 0.236, which is below break-even and produces the worst per-user result in the table.

Key findings and their implications

Finding 1 — a value function with a short horizon solves what the bandit could not. The difference between this notebook and the LinUCB notebook is not the features, the budget or the evaluation protocol, all of which are identical. It is that a Q-function propagates the delayed fatigue cost backwards to the action that caused it. Mean reward moves from -17.53 to +6.55 on exactly that change.

Finding 2 — selective silence is a learned behaviour and should be reported as a result, not repaired as a bug. The most valuable thing this agent learned is *which users not to contact*. A notification system that treats coverage as the objective would force sends onto NightShiftWorker and lose reward doing it; see the Double DQN notebook, where exactly that trade is measured.

Finding 3 — the discount factor is the binding hyperparameter. Section 11's sensitivity analysis shows the outcome is far more sensitive to γ than to the learning rate, the buffer size or the network width. γ decides how much of the fatigue tail the agent can see, and therefore whether contact is affordable in its own accounting at all.

Implication for deployment. The policy worth shipping is per-user, not global. Two of the five archetypes support aggressive contact, one supports almost none, and the remaining two need a more selective policy than this agent found. A single global send rate is the wrong control surface.

How to read a reward of `0.00`

`0.00` is not a missing value and not a crash. It is exactly what an agent earns by never sending: no clicks are collected, so no click reward is paid, and no send or fatigue cost is incurred either. Because three of the five archetypes score *negatively* under the Fixed-18:00 baseline, silence genuinely outscores the industry-standard schedule on those users. Any agent that discovers this is not malfunctioning -- it has found a real, and uncomfortable, local optimum.

The economics that govern every result here

A send is worth making only when its click probability clears the break-even click-through rate

$$\mathrm{CTR}_{\text{break-even}} = \frac{W_{\text{send}} + W_{\text{fat}} \cdot \kappa/(1-\lambda)}{R_{\text{click}}} = 0.340$$

The $\kappa/(1-\lambda)$ term is the whole story: because fatigue decays geometrically rather than instantly, one send commits the agent to a stream of future fatigue penalties. At $\gamma = 0.99$ over a 168-step episode that discounted tail is worth roughly a hundred steps of debt while the click pays once, which is why silence can dominate. Every finding above is a consequence of that single inequality.