

Context-Aware Notification Engine (CANE)

Double Deep Q-Network (Double DQN)

This notebook implements a **Double Deep Q-Network (Double DQN)** agent for the Context-Aware Notification Engine (CANE).

Double DQN addresses overestimation bias in standard DQN by decoupling action selection from value evaluation:

$$\text{Target} = r + \gamma \cdot Q_{\text{target}} \left(s', \underset{a}{\operatorname{argmax}} Q_{\text{online}}(s', a) \right) \cdot (1 - \text{done})$$

Notebook Structure

Section	Contents
1	Setup & Dependencies
2	Agent Interface
3	Environment Dynamics (CANEEnv)
4	Baselines & Evaluation Protocol
5	Environment Sanity Checks & Plotting Utilities
6	Double DQN Architecture & Agent Implementation
7	Correctness Verification Tests
8	Registry & Execution Driver
9	Learning Curves & Optimisation Diagnostics
10	Decision Distribution & Policy Behaviour
11	Hyperparameter Sensitivity Analysis
12	Statistical Robustness & Model Persistence

1. Setup & Dependencies

Import required libraries, define action spaces, observation feature layouts, and initialize execution parameters.

```
In [1]: import numpy as np
import pandas as pd
```

```

import matplotlib.pyplot as plt
from abc import ABC, abstractmethod
import torch
import torch.nn as nn

# --- Action Space -----
HOLD = 0      # Stay silent; fatigue decays
ENGAGE = 1    # Engagement notification
INCENTIVE = 2 # Incentive / promo notification

N_ACTIONS = 3
ACTION_NAMES = {HOLD: "Hold", ENGAGE: "Engage", INCENTIVE: "Incentive"}

# --- Observation Feature Indices -----
IDX_HOUR = 0      # Raw hour (0-23)
IDX_DAY = 1       # Raw day of week (0-6)
IDX_ACTIVE = 2    # In-app activity flag (0 or 1)
IDX_FATIGUE = 3   # Accumulated fatigue [0, 1]
IDX_RECENCY = 4   # Normalised time since last send [0, 1]

N_BLOCKS = 8      # 3-hour blocks per day
BLOCK_HOURS = 24 // N_BLOCKS

IDX_ACT_RATE = 5   # Block activity rate (5..12)
IDX_CLICK_RATE = 13 # Block click rate (13..20)
IDX_TYPE_CR = 21   # Click rate per message type (21..22)
IDX_SINCE_CLICK = 23 # Hours since last click

STATE_DIM_BASE = 5
STATE_DIM = 24

BELIEF_PRIOR = (1.0, 4.0)
USE_BELIEF = True

def agent_state_dim():
    return STATE_DIM if USE_BELIEF else STATE_DIM_BASE

DEVICE = "cpu"
torch.set_num_threads(1)

print("Actions:", ACTION_NAMES)
print("State dim:", STATE_DIM, "| agent sees:", agent_state_dim())
print("torch version:", torch.__version__, "| device:", DEVICE)

```

```

Actions: {0: 'Hold', 1: 'Engage', 2: 'Incentive'}
State dim: 24 | agent sees: 24
torch version: 2.13.0+cpu | device: cpu

```

2. Agent Interface

Abstract base class defining the shared execution contract for all agents.

```
In [2]: # --- The agent contract -----
# Every policy in this notebook -- the two non-learning baselines and the
# learner alike -- implements this interface, which is what lets the evaluation
# harness below treat them all identically. The comparison is only like-for-like
# because no agent gets a code path of its own.
#
# `act` returns (action, aux). The aux dict is how an agent hands the harness
# whatever it needs to record: the deep agents put their Q-values in it.
# `update` is a no-op for the baselines; `reset` is called at every episode
# boundary, which matters for any agent holding per-episode state.
class Agent(ABC):
    """Abstract base class for all CANE agents."""

    @property
    def name(self):
        return self.__class__.__name__

    @abstractmethod
    def act(self, state, greedy=False):
        """Select an action given an observation state."""
        raise NotImplementedError

    @abstractmethod
    def update(self, state, action, reward, next_state, done, aux):
        """Update policy / value function from a transition."""
        raise NotImplementedError

    def reset(self):
        """Reset agent state at episode boundaries (default no-op)."""
        pass

print("Agent interface defined.")
```

Agent interface defined.

3. Environment Dynamics

Defines user archetypes, click/fatigue/churn dynamics, and the `CANEnv` environment.

```
In [3]: # --- Environment Parameters -----
ARCHETYPE_W = {
    "OfficeWorker": {ENGAGE: 1.15, INCENTIVE: 1.05},
    "NightOwlStudent": {ENGAGE: 0.90, INCENTIVE: 1.60},
    "NightShiftWorker": {ENGAGE: 1.00, INCENTIVE: 1.25},
    "NormalStudent": {ENGAGE: 0.90, INCENTIVE: 1.55},
    "Housewife": {ENGAGE: 1.05, INCENTIVE: 1.45},
}

CANE_CONFIG = dict(
    steps_per_episode=168,
    lam=0.90,
    kappa={HOLD: 0.0, ENGAGE: 0.12, INCENTIVE: 0.18},
)
```

```

mu=0.80,
w_arch=ARCHETYPE_W,
churn_threshold=0.70,
gamma_0=-7.0,
gamma_1=4.0,
R_click=10.0,
W_send={HOLD: 0.0, ENGAGE: 1.0, INCENTIVE: 2.5},
W_fat=2.0,
R_churn=30.0,
beta=1.15,
streak_mode="marginal",
streak_cap=5.0,
streak_lapse_hours=36,
active_scale=0.6,
)

ARCHETYPES = ["OfficeWorker", "NightOwlStudent", "NightShiftWorker",
              "NormalStudent", "Housewife"]

def _bump(hours, centre, width, amp):
    delta = (hours - centre + 12.0) % 24.0 - 12.0
    return amp * np.exp(-0.5 * (delta / width) ** 2)

def archetype_curve(archetype, hours=None):
    h = np.arange(24.0) if hours is None else np.asarray(hours, dtype=float)
    if archetype == "OfficeWorker":
        c = _bump(h, 7.5, 1.0, 0.42) + _bump(h, 12.5, 1.1, 0.30) + _bump(h, 19.5, 2
    elif archetype == "NightOwlStudent":
        c = _bump(h, 0.0, 2.2, 0.32) + _bump(h, 22.0, 1.6, 0.26) + _bump(h, 16.0, 2
    elif archetype == "NightShiftWorker":
        c = _bump(h, 16.5, 1.6, 0.45) + _bump(h, 6.5, 1.5, 0.40) + _bump(h, 2.0, 1.
    elif archetype == "NormalStudent":
        c = _bump(h, 16.0, 1.6, 0.45) + _bump(h, 21.0, 1.5, 0.42) + _bump(h, 7.0, 1
    elif archetype == "Housewife":
        c = _bump(h, 10.0, 2.4, 0.42) + _bump(h, 14.5, 2.4, 0.40) + _bump(h, 20.5,
    else:
        raise ValueError(f"unknown archetype: {archetype}")
    return np.clip(c, 0.0, 1.0)

ARCHETYPE_TABLE = {a: archetype_curve(a) for a in ARCHETYPES}
WEEKEND_FLATTEN = {"OfficeWorker": 0.55, "NormalStudent": 0.50,
                  "NightShiftWorker": 0.30, "NightOwlStudent": 0.15,
                  "Housewife": 0.05}

print(f"{len(ARCHETYPES)} archetypes defined:", ", ".join(ARCHETYPES))

```

5 archetypes defined: OfficeWorker, NightOwlStudent, NightShiftWorker, NormalStudent, Housewife

```

In [4]: # --- The simulated user -----
# The environment is the assignment's model of one app user over a week: 168
# hourly steps, each of which is a decision point. There is no external dataset;
# this class IS the data-generating process, and the archetype tables above are
# its ground truth.
#
# The agent never observes the archetype or the true receptiveness curve. It

```

```

# sees only the belief state assembled in `_observe`, so learning who this
# person is has to happen through reward alone. That is the whole problem.
class CANEnv:
    """Context-Aware Notification Engine simulator environment."""

    def __init__(self, config=None, archetype=None, seed=0):
        self.cfg = {**CAN_CONFIG, **(config or {})}
        self.fixed_archetype = archetype
        self.rng = np.random.default_rng(seed)
        self.reset()

    def _p0(self):
        # Baseline receptiveness for this person at this hour of the day.
        # On weekends (day >= 5) it is blended toward the person's daily mean:
        # a weekday commuter's sharp 08:00 peak flattens out on a Saturday.
        base = ARCHETYPE_TABLE[self.archetype][self.hour]
        if self.day >= 5:
            f = WEEKEND_FLATTEN[self.archetype]
            base = (1.0 - f) * base + f * ARCHETYPE_TABLE[self.archetype].mean()
        return float(base)

    def _observe(self):
        # The belief state. Note what is absent: the archetype, the true
        # receptiveness curve, and the click probability. The agent gets only
        # what an app could actually log about a user.
        s = np.zeros(STATE_DIM, dtype=np.float32)
        s[IDX_HOUR] = self.hour
        s[IDX_DAY] = self.day
        s[IDX_ACTIVE] = self.active
        s[IDX_FATIGUE] = self.fatigue
        s[IDX_RECENCY] = self.recency
        # Beta-prior smoothing. Early in an episode a block may have zero sends,
        # and a raw clicks/sends ratio would be 0/0. Adding the prior makes an
        # unobserved block read as 'unknown' rather than as 'known to be dead'.
        a0, b0 = BELIEF_PRIOR
        s[IDX_ACT_RATE:IDX_ACT_RATE + N_BLOCKS] = (self.blk_act_hits + a0) / (self.
        s[IDX_CLICK_RATE:IDX_CLICK_RATE + N_BLOCKS] = (self.blk_clicks + a0) / (sel
        s[IDX_TYPE_CR] = (self.type_clicks[ENGAGE] + a0) / (self.type_sends[ENGAGE]
        s[IDX_TYPE_CR + 1] = (self.type_clicks[INCENTIVE] + a0) / (self.type_sends[
        s[IDX_SINCE_CLICK] = min(self.hours_since_click / 48.0, 1.0)
        return s

    def _roll_activity(self):
        # Whether the user is already in the app this hour. Rolled every step,
        # including on Hold, because the app observes app-opens whether or not it
        # sent anything -- that is the free signal the belief state accumulates.
        p = min(1.0, self.cfg["active_scale"] * self._p0())
        self.active = int(self.rng.random() < p)
        b = self.hour // BLOCK_HOURS
        self.blk_act_obs[b] += 1
        self.blk_act_hits[b] += self.active

    def reset(self, seed=None):
        if seed is not None:
            self.rng = np.random.default_rng(seed)
        # Episodes start at a random hour and day, so the agent cannot learn a

```

```

# fixed step-index-to-hour mapping and must read the hour from the state.
self.archetype = self.fixed_archetype or ARCHETYPES[self.rng.integers(len(A
self.t = 0
self.hour = int(self.rng.integers(24))
self.day = int(self.rng.integers(7))
self.fatigue = 0.0
self.rency = 1.0
self.hours_since_send = 24
self.streak = 0
self.hours_since_click = 0
self.opted_out = False
self.blk_act_obs = np.zeros(N_BLOCKS)
self.blk_act_hits = np.zeros(N_BLOCKS)
self.blk_sends = np.zeros(N_BLOCKS)
self.blk_clicks = np.zeros(N_BLOCKS)
self.type_sends = np.zeros(N_ACTIONS)
self.type_clicks = np.zeros(N_ACTIONS)
self._roll_activity()
return self._observe(), {"archetype": self.archetype}

def step(self, action):
    cfg = self.cfg
    action = int(action)
    sent = action != HOLD
    f_now = self.fatigue

    clicked = False
    # A click is only possible on a send, and only when the user is NOT
    # already active in the app: a notification delivered to someone already
    # using the app earns nothing. Click probability then scales the hour's
    # receptiveness down by current fatigue and by the per-archetype weight
    # of this message type.
    if sent and not self.active:
        w_a = cfg["w_arch"][self.archetype][action]
        p_click = np.clip(self._p0() * (1.0 - cfg["mu"] * f_now) * w_a, 0.0, 1.
        clicked = bool(self.rng.random() < p_click)

    blk = self.hour // BLOCK_HOURS
    if sent:
        self.blk_sends[blk] += 1
        self.blk_clicks[blk] += clicked
        self.type_sends[action] += 1
        self.type_clicks[action] += clicked

    if clicked:
        self.streak += 1
        self.hours_since_click = 0
    else:
        self.hours_since_click += 1
        if self.hours_since_click > cfg["streak_lapse_hours"]:
            self.streak = 0

    # Consecutive clicks compound, which is what makes this a sequential
    # problem rather than a per-hour one: the value of a click depends on
    # whether the previous ones landed.
    streak_bonus = 0.0

```

```

if clicked and self.streak > 0:
    if cfg["streak_mode"] == "cumulative":
        streak_bonus = sum(cfg["beta"] ** k for k in range(1, self.streak + 1))
    else:
        streak_bonus = cfg["beta"] ** self.streak
    if cfg["streak_cap"] is not None:
        streak_bonus = min(streak_bonus, cfg["streak_cap"])

# Fatigue decays geometrically at `lam` and takes on `kappa[action]` per
# send. Because lam < 1, one send's cost is paid back over many future
# steps -- this recursion is the reason a send is expensive and the
# reason a patient agent can rationally choose to send nothing at all.
self.fatigue = float(np.clip(cfg["lam"] * f_now + cfg["kappa"][action], 0.0, 1.0))

# Past the threshold the user may opt out permanently, ending the
# episode. This is the hard ceiling on over-sending: no recovery.
if self.fatigue > cfg["churn_threshold"]:
    hazard = 1.0 / (1.0 + np.exp(-(cfg["gamma_0"] + cfg["gamma_1"] * self.fatigue)))
    self.opted_out = bool(self.rng.random() < hazard)

# Reward = click payoff, minus the cost of sending, minus a penalty on
# fatigue, minus the churn charge, plus any streak bonus.
# Fatigue is charged at `f_now`, the level BEFORE this step's send, so a
# send is not billed for the fatigue it has only just created -- that
# arrives on the following steps via the recursion above.
reward = (cfg["R_click"] * clicked
          - cfg["W_send"][action]
          - cfg["W_fat"] * f_now
          - cfg["R_churn"] * self.opted_out
          + streak_bonus)

self.t += 1
self.hour = (self.hour + 1) % 24
if self.hour == 0:
    self.day = (self.day + 1) % 7

# Recency, capped at 24h and normalised, is the agent's only direct view
# of how recently it last made contact.
self.hours_since_send = 0 if sent else self.hours_since_send + 1
self.recency = float(min(self.hours_since_send / 24.0, 1.0))
self._roll_activity()

terminated = self.opted_out
truncated = self.t >= cfg["steps_per_episode"]

info = {
    "archetype": self.archetype,
    "clicked": clicked, "sent": sent,
    "fatigue": self.fatigue, "fatigue_pre": f_now,
    "opted_out": self.opted_out, "streak": self.streak,
    "streak_bonus": streak_bonus,
    "hour": int((self.hour - 1) % 24), "active_pre": self.active,
}
return self._observe(), float(reward), terminated, truncated, info

print("CANEnv defined.")

```

4. Baselines & Evaluation Protocol

Provides benchmark agents (`FixedScheduleAgent` , `RandomAgent`) and evaluation harness functions.

```
In [5]: # --- Non-Learning baselines -----
# Fixed-18:00 is the industry default this project argues against: one
# hard-coded send time for everyone. Random is the lower bound. Neither
# learns, so both are evaluated once rather than across training seeds.
class FixedScheduleAgent(Agent):
    """Sends an Engagement Nudge at a fixed hour every day."""
    def __init__(self, hour=18, action=ENGAGE):
        self.hour, self.action = hour, action

    @property
    def name(self):
        return f"Fixed-{{self.hour:02d}}:00"

    def act(self, state, greedy=False):
        return (self.action if int(state[IDX_HOUR]) == self.hour else HOLD), {}

    def update(self, *args, **kwargs):
        return {}

class RandomAgent(Agent):
    """Selects actions uniformly at random."""
    def __init__(self, seed=0):
        self.rng = np.random.default_rng(seed)

    @property
    def name(self):
        return "Random"

    def act(self, state, greedy=False):
        return int(self.rng.integers(N_ACTIONS)), {}

    def update(self, *args, **kwargs):
        return {}

# --- The evaluation harness -----
# One function drives both training and evaluation, which is what keeps
# them honest: `learn=True, greedy=False` trains, `learn=False,
# greedy=True` scores a frozen policy. Passing explicit `seeds` replays
# an identical set of episodes for every agent.
def run_episodes(agent, env, n_episodes=None, seeds=None, learn=True,
                 greedy=False, collect_log=False):
    """Runs agent in env over specified episodes and returns evaluation metrics."""
    if seeds is None:
```



```

seeds = [None] * int(n_episodes)

ep_rewards, ep_sends, ep_clicks, ep_optouts, ep_lengths = [], [], [], [], []
# Counts actions by clock hour, which is what the policy heatmaps and the
# personalisation analysis are built from later.
hour_actions = np.zeros((24, N_ACTIONS), dtype=int)
log = []

for ep, sd in enumerate(seeds):
    state, info = env.reset(seed=sd)
    agent.reset()
    total_r = sends = clicks = 0.0

    while True:
        action, aux = agent.act(state, greedy=greedy)
        next_state, reward, terminated, truncated, info = env.step(action)

        if learn:
            agent.update(state, action, reward, next_state, terminated, aux)

        total_r += reward
        sends += info["sent"]
        clicks += info["clicked"]
        hour_actions[info["hour"], action] += 1

        if collect_log:
            log.append({"episode": ep, "step": env.t, "hour": info["hour"],
                       "archetype": info["archetype"], "action": action,
                       "clicked": info["clicked"], "sent": info["sent"],
                       "fatigue": info["fatigue"], "reward": reward,
                       "streak": info["streak"], "opted_out": info["opted_out"]})

        state = next_state
        if terminated or truncated:
            ep_optouts.append(float(terminated))
            break

    ep_rewards.append(total_r)
    ep_sends.append(sends)
    ep_clicks.append(clicks)
    ep_lengths.append(env.t)

# CTR is clicks per SEND, not per step, and is defined as 0.0 when the
# agent never sent. A never-send policy therefore reports ctr 0.00 and
# reward 0.00 -- neither is a missing value; both are what silence earns.
total_sends = float(np.sum(ep_sends))
metrics = {
    "agent": agent.name,
    "reward_mean": float(np.mean(ep_rewards)),
    "reward_std": float(np.std(ep_rewards)),
    "ctr": float(np.sum(ep_clicks) / total_sends) if total_sends > 0 else 0.0,
    "sends_per_episode": float(np.mean(ep_sends)),
    "clicks_per_episode": float(np.mean(ep_clicks)),
    "optout_rate": float(np.mean(ep_optouts)),
    "episode_length": float(np.mean(ep_lengths)),
}

```

```

    return metrics, (log if collect_log else None), np.array(ep_rewards), hour_acti

# --- The shared evaluation protocol -----
# These four values are identical in all four notebooks. That identity is
# what makes the cross-algorithm comparison valid, and `check_results.py`
# verifies it afterwards by confirming the deterministic baseline rows
# agree exactly across the four result files.
#
# The 900,000+ evaluation seeds are disjoint from the 1000+ training and
# 7000+ evaluation environment seeds below, so nothing is scored on an
# episode it was trained on. QUICK is a smoke-test switch: its numbers
# are NOT reportable.
QUICK = False

EVAL_SEEDS = list(range(900_000, 900_050 if QUICK else 900_200))
TRAIN_EPISODES = 150 if QUICK else 600
N_SEEDS = 2 if QUICK else 5

print(f"Eval protocol: {len(EVAL_SEEDS)} reserved episodes; "
      f"training budget {TRAIN_EPISODES} episodes; {N_SEEDS} seeds."
      + (" [QUICK]" if QUICK else ""))

```

Eval protocol: 200 reserved episodes; training budget 600 episodes; 5 seeds.

```

In [6]: # --- Frozen-policy probes -----
# Snapshots taken during training, used for the learning curves and the
# animation. A separate seed block again, so probing never touches the
# episodes used for training or for the reported score.
PROBE_SEEDS = list(range(950_000, 950_020))
SNAPSHOT_ARCHETYPES = ("OfficeWorker", "NightOwlStudent")

def policy_snapshot(agent, archetypes=SNAPSHOT_ARCHETYPES, seeds=PROBE_SEEDS):
    """Evaluate frozen greedy policy across probe seeds."""
    out = {}
    for arch in archetypes:
        env = CANEnv(seed=8000, archetype=arch)
        m, _, _, hours = run_episodes(agent, env, seeds=seeds, learn=False, greedy=
        out[arch] = {"hour_actions": hours, "reward": m["reward_mean"],
                    "ctr": m["ctr"], "sends": m["sends_per_episode"],
                    "optout": m["optout_rate"]}

    return out

# Trains a fresh agent per seed and scores each on the held-out episodes.
# A fresh agent per seed, not one agent re-evaluated: the spread being
# reported is the variance of the LEARNING process, not of the scoring.
def train_and_evaluate(make_agent, n_seeds=None, train_episodes=TRAIN_EPISODES,
                       archetype=None, collect_hours=False,
                       snapshot_every=None, snapshot_seed=0,
                       snapshot_archetypes=SNAPSHOT_ARCHETYPES):
    """Train an agent across multiple random seeds and evaluate on held-out seeds."
    n_seeds = N_SEEDS if n_seeds is None else n_seeds
    per_seed, hour_stack, snapshots = [], [], []

    for seed in range(n_seeds):
        agent = make_agent(seed)

```

```

# Distinct environment seeds for training and evaluation, so a policy
# cannot succeed by memorising particular episodes.
train_env = CANEnv(seed=1000 + seed, archetype=archetype)

if snapshot_every and seed == snapshot_seed:
    snapshots.append((0, policy_snapshot(agent, snapshot_archetypes)))
    done = 0
    while done < train_episodes:
        chunk = min(snapshot_every, train_episodes - done)
        run_episodes(agent, train_env, n_episodes=chunk, learn=True, greedy
        done += chunk
        snapshots.append((done, policy_snapshot(agent, snapshot_archetypes)
    else:
        run_episodes(agent, train_env, n_episodes=train_episodes, learn=True, g

eval_env = CANEnv(seed=7000 + seed, archetype=archetype)
m, _, rewards, hours = run_episodes(agent, eval_env, seeds=EVAL_SEEDS, learn
per_seed.append(m)
hour_stack.append(hours)

rewards = np.array([m["reward_mean"] for m in per_seed])
summary = {
    "agent": per_seed[0]["agent"],
    "reward_mean": rewards.mean(),
    "reward_std": rewards.std(ddof=0),
    # ddof=1: the sample standard deviation across seeds, which is what the
    # confidence intervals in the figures are derived from.
    "reward_sem_std": rewards.std(ddof=1) if n_seeds > 1 else 0.0,
    "ctr": np.mean([m["ctr"] for m in per_seed]),
    "sends_per_episode": np.mean([m["sends_per_episode"] for m in per_seed]),
    "optout_rate": np.mean([m["optout_rate"] for m in per_seed]),
    "per_seed_rewards": rewards,
}
if collect_hours:
    summary["hour_actions"] = np.sum(hour_stack, axis=0)
if snapshots:
    summary["snapshots"] = snapshots
return summary

# Baselines have nothing to train, so they are evaluated once. Both are
# deterministic given the seed list, which is precisely why their rows
# can be used to prove the four notebooks share one protocol.
def evaluate_baseline(make_agent, archetype=None):
    agent = make_agent(0)
    env = CANEnv(seed=7000, archetype=archetype)
    m, _, rewards, hours = run_episodes(agent, env, seeds=EVAL_SEEDS, learn=False,
    return {"agent": m["agent"], "reward_mean": m["reward_mean"],
            "reward_std": m["reward_std"], "reward_sem_std": 0.0,
            "ctr": m["ctr"], "sends_per_episode": m["sends_per_episode"],
            "optout_rate": m["optout_rate"],
            "per_seed_rewards": np.array([m["reward_mean"]]),
            "hour_actions": hours}

```

5. Environment Sanity Checks & Exploratory Data Analysis (EDA)

Validates environment dynamics and visualises baseline responsiveness curves, fatigue dynamics, click surfaces, and ground-truth policy reference maps.

```
In [7]: # --- 5.1 Environment Sanity Checks -----
env = CANEnv(seed=0)

e = CANEnv(archetype="Housewife", seed=1)
e.reset(seed=1)
f_trace, expected = [], 0.0
for _ in range(30):
    _, _, term, trunc, info = e.step(ENGAGE)
    expected = min(0.9 * expected + CANE_CONFIG["kappa"][ENGAGE], 1.0)
    f_trace.append((info["fatigue"], expected))
    if term or trunc:
        break
assert all(abs(a - b) < 1e-9 for a, b in f_trace)

e.reset(seed=2)
for _ in range(5):
    e.step(ENGAGE)
f_high = e.fatigue
for _ in range(10):
    e.step(HOLD)
assert e.fatigue < f_high

class _AlwaysHold(Agent):
    def act(self, s, greedy=False): return HOLD, {}
    def update(self, *a, **k): return {}

m_hold, _, _, _ = run_episodes(_AlwaysHold(), CANEnv(seed=3), 30, learn=False, greedy=True)
assert m_hold["optout_rate"] == 0.0

class _AlwaysSend(Agent):
    def act(self, s, greedy=False): return ENGAGE, {}
    def update(self, *a, **k): return {}

m_spam, _, _, _ = run_episodes(_AlwaysSend(), CANEnv(seed=4), 30, learn=False, greedy=True)
assert m_spam["optout_rate"] > 0.8

print("Environment sanity checks passed.")
```

Environment sanity checks passed.

```
In [8]: # --- Plot G1: Receptive Windows by Archetype ---
from pathlib import Path

FOCUS = ["OfficeWorker", "NightOwlStudent"]
ARCH_COLOUR = {"OfficeWorker": "#2f6fb5", "NightOwlStudent": "#c4453c"}

FIGDIR = Path("figures")
FIGDIR.mkdir(exist_ok=True)
```


```

def save(fig, tag):
    fig.savefig(FIGDIR / f"{tag}.png", dpi=150, bbox_inches="tight")

def weekend_curve(archetype):
    base = ARCHETYPE_TABLE[archetype]
    f = WEEKEND_FLATTEN[archetype]
    return (1.0 - f) * base + f * base.mean()

hours = np.arange(24)
fig, axes = plt.subplots(1, 2, figsize=(10.5, 3.4), sharey=True)
for ax, arch in zip(axes, FOCUS):
    ax.fill_between(hours, ARCHETYPE_TABLE[arch], color=ARCH_COLOUR[arch], alpha=0.
    ax.plot(hours, ARCHETYPE_TABLE[arch], color=ARCH_COLOUR[arch], lw=2.0, label="w
    ax.plot(hours, weekend_curve(arch), color=ARCH_COLOUR[arch], lw=1.3, ls="--", l
    ax.set_title(arch, fontsize=10)
    ax.set_xlabel("hour of day")
    ax.set_xticks(range(0, 24, 3))
    ax.legend(frameon=False, fontsize=8)
axes[0].set_ylabel(r"$p_0$ click propensity")
fig.suptitle("G1 Receptive windows by archetype", fontsize=11)
fig.tight_layout()
save(fig, "G1_archetype_curves")
plt.show()

```



```

In [9]: # --- Plot G2: Fatigue Dynamics ---
lam = CANE_CONFIG["lam"]
kap = CANE_CONFIG["kappa"]
thr = CANE_CONFIG["churn_threshold"]
pattern = np.array([1, 1, 1, 1, 1, 1] + [0] * 18)

fig, ax = plt.subplots(figsize=(9, 3.0))
ax.fill_between(range(len(pattern)), 0, 1, where=pattern.astype(bool),
                color="#999999", alpha=0.12, step="post")
for act, col in [(ENGAGE, "#2f6fb5"), (INCENTIVE, "#c4453c")]:
    F, trace = 0.0, []
    for a in pattern:
        F = float(np.clip(lam * F + (kap[act] if a else 0.0), 0.0, 1.0))
        trace.append(F)
    ax.step(range(len(trace)), trace, where="post", lw=1.8, color=col,
            label=f"{ACTION_NAMES[act]} (kappa={kap[act]})")


ax.axhline(thr, color="#666666", ls=":", lw=1.4)

```

```

ax.text(len(pattern) - 0.5, thr + 0.015, f"churn threshold {thr}",
       ha="right", va="bottom", fontsize=8, color="#666666")
ax.set_xlabel("hours")
ax.set_ylabel("fatigue $F$")
ax.set_ylim(0, 1)
ax.legend(frameon=False, fontsize=8, loc="upper right")
ax.set_title("G2 Six hourly sends (shaded), then silence", fontsize=11)
fig.tight_layout()
save(fig, "G2_fatigue_dynamics")
plt.show()

```




```

In [10]: # --- Plot G3: Click Probability Surface ---
mu = CANE_CONFIG["mu"]

def break_even(archetype, action):
    cfg = CANE_CONFIG
    debt = cfg["W_fat"] * cfg["kappa"][action] / (1.0 - cfg["lam"])
    return (debt + cfg["W_send"][action]) / (cfg["R_click"] * cfg["w_arch"][archetype])

F_grid = np.linspace(0.0, 1.0, 101)
fig, axes = plt.subplots(1, 2, figsize=(11, 3.5), sharey=True)
for ax, arch in zip(axes, FOCUS):
    P = np.clip(np.outer(1.0 - mu * F_grid, ARCHETYPE_TABLE[arch])
               * CANE_CONFIG["w_arch"][arch][ENGAGE], 0.0, 1.0)
    im = ax.imshow(P, origin="lower", aspect="auto", cmap="magma",
                  extent=[0, 23, 0, 1], vmin=0, vmax=0.6)
    cs = ax.contour(np.arange(24), F_grid, P, levels=[break_even(arch, ENGAGE)],
                   colors="w", linewidths=1.5)
    ax.clabel(cs, fmt={break_even(arch, ENGAGE): "break-even"}, fontsize=7)
    ax.set_title(f"{arch} (Engage break-even {break_even(arch, ENGAGE):.2f})", fontweight="bold")
    ax.set_xlabel("hour of day")
    ax.set_xticks(range(0, 24, 3))
axes[0].set_ylabel("fatigue $F$")
fig.colorbar(im, ax=axes, label="P(click | Engage)")
fig.suptitle("G3 Engage click probability surface", fontsize=11)
save(fig, "G3_click_surface")
plt.show()

```




```
In [11]: # --- Plot G4: Ground-Truth Policy Reference ---
def net_value(archetype, action, hour, F):
    cfg = CANE_CONFIG
    p = min(ARCHETYPE_TABLE[archetype][hour] * (1.0 - cfg["mu"] * F)
            * cfg["w_arch"][archetype][action], 1.0)
    debt = cfg["W_fat"] * cfg["kappa"][action] / (1.0 - cfg["lam"])
    return p * cfg["R_click"] - cfg["W_send"][action] - debt

F_GRID = np.linspace(0.0, 1.0, 101)

GROUND_TRUTH = {}
for arch in ARCHETYPES:
    M = np.zeros((len(F_GRID), 24), dtype=int)
    for j, F in enumerate(F_GRID):
        for h in range(24):
            v = {a: net_value(arch, a, h, F) for a in (ENGAGE, INCENTIVE)}
            best = max(v, key=v.get)
            M[j, h] = best if v[best] > 0 else HOLD
    GROUND_TRUTH[arch] = M

from matplotlib.colors import ListedColormap
ACTION_COLOUR = ["#e8e8e8", "#2f6fb5", "#c4453c"]

fig, axes = plt.subplots(1, 2, figsize=(11, 3.5), sharey=True)
for ax, arch in zip(axes, FOCUS):
    ax.imshow(GROUND_TRUTH[arch], origin="lower", aspect="auto",
              cmap=ListedColormap(ACTION_COLOUR), vmin=0, vmax=2, extent=[0, 23, 0,
    ax.set_title(arch, fontsize=10)
    ax.set_xlabel("hour of day")
    ax.set_xticks(range(0, 24, 3))
axes[0].set_ylabel("fatigue $F$")
axes[1].legend([plt.Rectangle((0, 0), 1, 1, color=col) for col in ACTION_COLOUR],
               ["Hold", "Engage", "Incentive"], frameon=False, fontsize=8, loc="upp
fig.suptitle("G4 Myopic-optimal action by hour and fatigue", fontsize=11)
fig.tight_layout()
save(fig, "G4_ground_truth_policy")
plt.show()
```



6. Double DQN Agent Implementation

Defines `DDQN_CONFIG`, `ReplayBuffer`, `QNetwork`, and the `DDQNAgent` class.

```
In [12]: # --- Configuration -----
# Retuned after the first full run collapsed. The agent contacted 2 of the 5
# archetypes and scored exactly 0.00 on the other 3 -- which is not a missing
# value, it is precisely what an agent earns by never sending: no clicks
# collected, no send or fatigue costs paid.
#
# gamma is the lever. An episode is 168 steps, so at gamma=0.99 the discounted
# tail of a send's fatigue cost is worth on the order of 100 steps of future
# debt, while the click payoff arrives once and immediately. Under that
# accounting silence really is the optimal policy, and no amount of extra
# training escapes it. At gamma=0.90 the fatigue tail is worth roughly 10 steps
# and contact becomes affordable again.
#
# The other four changes come with it: a faster learning rate and a larger
# buffer to exploit the richer signal, plus a slightly higher exploration floor
# held over a longer decay so the agent keeps trying sends after the early ones
# look expensive. These are the values the DQN notebook already uses, so the two
# value-based agents are now tuned alike -- which is what keeps the comparison
# between them about the algorithm rather than about the hyperparameters.
DDQN_CONFIG = dict(
    lr=3e-4,           # Adam Learning rate           (was 1e-4)
    batch_size=64,     # Minibatch size
    max_grad_norm=0.5, # Gradient clipping norm
    gamma=0.90,        # Discount factor           (was 0.99)
    epsilon_start=1.0, # Initial exploration rate
    epsilon_end=0.08,  # Final exploration rate    (was 0.05)
    epsilon_decay_steps=60000, # Linear decay duration, steps (was 50000)
    buffer_size=20000, # Replay buffer capacity    (was 10000)
    min_replay_size=1000, # Warmup steps before Learning
    target_update_freq=500, # Target network hard update frequency
    hidden_sizes=(64, 64), # Q-network hidden layer dimensions
)

print("DDQN_CONFIG:", ", ".join(f"{k}={v}" for k, v in DDQN_CONFIG.items()))
```


DDQN_CONFIG: lr=0.0003, batch_size=64, max_grad_norm=0.5, gamma=0.9, epsilon_start=1.0, epsilon_end=0.08, epsilon_decay_steps=60000, buffer_size=20000, min_replay_size=1000, target_update_freq=500, hidden_sizes=(64, 64)

```
In [13]: # --- Experience replay -----
# Off-policy Learning needs decorrelated samples: consecutive hours in this
# environment are highly correlated (fatigue decays smoothly, the hour advances
# by one), and training on them in order would violate the i.i.d. assumption the
# gradient step relies on and make the updates unstable.
#
# Preallocated NumPy arrays with a circular write pointer, rather than a deque
# of tuples: fixed memory, and sampling is one fancy-index instead of a Python
# loop over `batch_size` objects.
class ReplayBuffer:
    """Experience replay buffer storing (state, action, reward, next_state, done)."""

    def __init__(self, capacity, state_dim, seed=0):
        self.capacity = capacity
        self.rng = np.random.default_rng(seed)
        self.pos = 0
        self.size = 0

        # Preallocated to `capacity` up front; `pos` wraps, so the buffer holds a
        # sliding window of the most recent transitions.
        self.states = np.zeros((capacity, state_dim), dtype=np.float32)
        self.actions = np.zeros(capacity, dtype=np.int64)
        self.rewards = np.zeros(capacity, dtype=np.float32)
        self.next_states = np.zeros((capacity, state_dim), dtype=np.float32)
        self.dones = np.zeros(capacity, dtype=np.float32)

    def push(self, state, action, reward, next_state, done):
        self.states[self.pos] = state
        self.actions[self.pos] = action
        self.rewards[self.pos] = reward
        self.next_states[self.pos] = next_state
        self.dones[self.pos] = float(done)
        self.pos = (self.pos + 1) % self.capacity
        self.size = min(self.size + 1, self.capacity)

    def sample(self, batch_size):
        # Uniform sampling with replacement over the filled region only.
        idx = self.rng.integers(0, self.size, size=batch_size)
        return (
            torch.from_numpy(self.states[idx]),
            torch.from_numpy(self.actions[idx]),
            torch.from_numpy(self.rewards[idx]),
            torch.from_numpy(self.next_states[idx]),
            torch.from_numpy(self.dones[idx]),
        )

    def __len__(self):
        return self.size

print("ReplayBuffer defined.")
```

ReplayBuffer defined.

```
In [14]: # --- The Q-network -----
# A small MLP mapping the 24-dimensional belief state to one Q-value per action.
# Deliberately small: the state is low-dimensional and the training budget is
# 600 episodes, so capacity is not the binding constraint here and a larger net
# would only overfit the replay buffer faster.
class QNetwork(nn.Module):
    """Multi-Layer Perceptron Q-network."""

    def __init__(self, d_in, n_actions, hidden_sizes):
        super().__init__()
        self.net = self._mlp(d_in, hidden_sizes, n_actions, out_gain=0.01)

    @staticmethod
    def _mlp(d_in, hidden, d_out, out_gain):
        # Orthogonal initialisation with gain sqrt(2) suits the Tanh activations;
        # the output head uses a much smaller gain (0.01) so every Q-value starts
        # near zero and no action is arbitrarily favoured before any learning.
        layers, prev = [], d_in
        for h in hidden:
            lin = nn.Linear(prev, h)
            nn.init.orthogonal_(lin.weight, gain=np.sqrt(2.0))
            nn.init.zeros_(lin.bias)
            layers += [lin, nn.Tanh()]
            prev = h
        head = nn.Linear(prev, d_out)
        nn.init.orthogonal_(head.weight, gain=out_gain)
        nn.init.zeros_(head.bias)
        layers.append(head)
        return nn.Sequential(*layers)

    def forward(self, x):
        return self.net(x)

print("QNetwork defined.")
```

QNetwork defined.

```
In [15]: class DDQNAgent(Agent):
    """Double Deep Q-Network Agent."""

    def __init__(self, seed=0, label=None, **overrides):
        unknown = set(overrides) - set(DDQN_CONFIG)
        if unknown:
            raise ValueError(f"unknown DDQN hyperparameter(s): {sorted(unknown)}")
        self.cfg = {**DDQN_CONFIG, **overrides}
        self._label = label

        self.rng = np.random.default_rng(seed)
        torch.manual_seed(seed)
        self.seed = seed

        self.d_in = agent_state_dim()
        self.online_net = QNetwork(self.d_in, N_ACTIONS, tuple(self.cfg["hidden_sizes"]))
        self.target_net = QNetwork(self.d_in, N_ACTIONS, tuple(self.cfg["hidden_sizes"]))
        self._sync_target()
        self.target_net.eval()
```

```

self.opt = torch.optim.Adam(self.online_net.parameters(), lr=self.cfg["lr"])
self.buffer = ReplayBuffer(self.cfg["buffer_size"], self.d_in, seed=seed)

self.history = []
self.n_updates = 0
self.total_steps = 0

def _features(self, state):
    x = np.asarray(state, dtype=np.float32)[:self.d_in].copy()
    x[IDX_HOUR] = x[IDX_HOUR] / 23.0
    x[IDX_DAY] = x[IDX_DAY] / 6.0
    return x

def _epsilon(self):
    cfg = self.cfg
    frac = min(1.0, self.total_steps / max(1, cfg["epsilon_decay_steps"]))
    return cfg["epsilon_start"] + frac * (cfg["epsilon_end"] - cfg["epsilon_start"])

def _sync_target(self):
    self.target_net.load_state_dict(self.online_net.state_dict())

@property
def name(self):
    return self._label or "DDQN"

def act(self, state, greedy=False):
    x = torch.from_numpy(self._features(state)).unsqueeze(0)
    with torch.no_grad():
        q_values = self.online_net(x).squeeze(0)

    if greedy:
        action = int(torch.argmax(q_values).item())
    else:
        eps = self._epsilon()
        if self.rng.random() < eps:
            action = int(self.rng.integers(N_ACTIONS))
        else:
            action = int(torch.argmax(q_values).item())

    return action, {"q_values": q_values.numpy().tolist()}

def update(self, state, action, reward, next_state, done, aux):
    feat_s = self._features(state)
    feat_ns = self._features(next_state)

    self.buffer.push(feat_s, action, reward, feat_ns, done)
    self.total_steps += 1

    if len(self.buffer) < self.cfg["min_replay_size"]:
        return {}

    diag = self._optimise()

    if self.total_steps % self.cfg["target_update_freq"] == 0:
        self._sync_target()

```

```

        return diag

def reset(self):
    pass

def _optimise(self):
    cfg = self.cfg
    states, actions, rewards, next_states, dones = self.buffer.sample(cfg["batch_size"])

    q_all = self.online_net(states)
    q_values = q_all.gather(1, actions.unsqueeze(1).long()).squeeze(1)

    # Double DQN target decoupling
    with torch.no_grad():
        next_q_online = self.online_net(next_states)
        best_actions = next_q_online.argmax(dim=1, keepdim=True)
        next_q_target = self.target_net(next_states)
        next_q = next_q_target.gather(1, best_actions).squeeze(1)
        target = rewards + cfg["gamma"] * next_q * (1.0 - dones)

    loss = nn.functional.mse_loss(q_values, target)

    self.opt.zero_grad()
    loss.backward()
    nn.utils.clip_grad_norm_(self.online_net.parameters(), cfg["max_grad_norm"])
    self.opt.step()

    self.n_updates += 1
    diag = {
        "loss": float(loss.detach()),
        "mean_q": float(q_values.mean().detach()),
        "max_q": float(q_values.max().detach()),
        "epsilon": self._epsilon(),
        "update": self.n_updates,
        "step": self.total_steps,
    }
    self.history.append(diag)
    return diag

def state_dict_for_save(self):
    return {"online_net": self.online_net.state_dict(),
            "target_net": self.target_net.state_dict(),
            "cfg": dict(self.cfg),
            "d_in": self.d_in, "seed": self.seed,
            "n_updates": self.n_updates, "total_steps": self.total_steps}

_probe = DDQNAgent(seed=0)
print("DDQNAgent defined:", _probe)
print("parameters (online):", sum(p.numel() for p in _probe.online_net.parameters()))

```

DDQNAgent defined: <__main__.DDQNAgent object at 0x0000020896975D90>
parameters (online): 5955

7. Correctness Verification Tests

Suite of automated unit tests ensuring interface adherence, target computation logic, evaluation purity, known-optimum convergence, and seed reproducibility.

```
In [16]: # --- 7.1 Interface Conformance -----
_env = CANEEEnv(seed=0)
_s0, _ = _env.reset(seed=123)

assert isinstance(_probe, Agent)
assert _probe.name == "DDQN"

_a, _aux = _probe.act(_s0, greedy=False)
assert isinstance(_a, int) and _a in (HOLD, ENGAGE, INCENTIVE)
assert isinstance(_aux, dict)

_greedy = {_probe.act(_s0, greedy=True)[0] for _ in range(100)}
assert len(_greedy) == 1

_s1, _r, _term, _trunc, _ = _env.step(_a)
_out = _probe.update(_s0, _a, _r, _s1, _term, _aux)
assert isinstance(_out, dict)

print("7.1 interface conformance PASS")
```

7.1 interface conformance PASS

```

In [17]: # --- 7.2 Double DQN Target Logic Test -----
_test_agent = DDQNAgent(seed=42, min_replay_size=10, target_update_freq=1000)
_test_env = CANEEEnv(seed=999)

for _ in range(5):
    _s, _ = _test_env.reset(seed=None)
    for _ in range(168):
        _a, _aux = _test_agent.act(_s, greedy=False)
        _ns, _r, _term, _trunc, _ = _test_env.step(_a)
        _test_agent.update(_s, _a, _r, _ns, _term, _aux)
        _s = _ns
        if _term or _trunc:
            break

_test_state = torch.randn(1, _test_agent.d_in)
with torch.no_grad():
    _q_online = _test_agent.online_net(_test_state)
    _best_a = _q_online.argmax(dim=1)
    _q_target = _test_agent.target_net(_test_state)
    _target_val = _q_target[0, _best_a[0]]
    _vanilla_val = _q_target.max(dim=1).values[0]

_online_params = list(_test_agent.online_net.parameters())
_target_params = list(_test_agent.target_net.parameters())
_any_different = any(not torch.equal(o, t) for o, t in zip(_online_params, _target_
assert _any_different

print("7.2 Double DQN target logic PASS")

```

7.2 Double DQN target logic PASS

```

In [18]: # --- 7.3 Evaluation Purity Test -----
_check = DDQNAgent(seed=3)
_train_env = CANEEEnv(seed=4321)
run_episodes(_check, _train_env, n_episodes=10, learn=True, greedy=False)
_buf_size_before = len(_check.buffer)

_before = [p.detach().clone() for p in _check.online_net.parameters()]
_updates_before = _check.n_updates

_m, _, _, _ = run_episodes(_check, CANEEEnv(seed=4321), n_episodes=5, learn=False, g
_after = [p.detach().clone() for p in _check.online_net.parameters()]

assert all(torch.equal(a, b) for a, b in zip(_before, _after))
assert _check.n_updates == _updates_before
assert len(_check.buffer) == _buf_size_before

print("7.3 evaluation purity PASS")

```

7.3 evaluation purity PASS

```

In [19]: # --- 7.4 Known-Optimum Learning Test -----
def run_synthetic_ddqn(agent, theta_star, n_episodes=80, ep_len=168, seed=0):
    rng = np.random.default_rng(seed)
    d = theta_star.shape[1]
    hits, rewards = [], []

```

```

for _ in range(n_episodes):
    agent.reset()
    for t in range(ep_len):
        x = rng.normal(size=d).astype(np.float32)
        best = int(np.argmax(theta_star @ agent._features(x)))
        action, aux = agent.act(x, greedy=False)
        reward = 1.0 if action == best else 0.0
        x_next = rng.normal(size=d).astype(np.float32)
        agent.update(x, action, reward, x_next, False, aux)
        hits.append(action == best)
        rewards.append(reward)
    return np.array(hits), np.array(rewards)

_syn_agent = DDQNAgent(seed=0, gamma=0.0, epsilon_decay_steps=5000)
_theta = np.random.default_rng(42).normal(size=(N_ACTIONS, _syn_agent.d_in))
_hits, _rew = run_synthetic_ddqn(_syn_agent, _theta, n_episodes=80)

_first = _hits[:len(_hits) // 4].mean()
_last = _hits[-len(_hits) // 4:].mean()

assert _last > 0.65
assert _last > _first + 0.10
print("7.4 known-optimum learning PASS")

```

7.4 known-optimum learning PASS

```

In [20]: # --- 7.5 Shared Environment Verification -----
assert len(ARCHETYPES) == 5
assert CANE_CONFIG["steps_per_episode"] == 168
print("7.5 shared environment verification PASS")

```

7.5 shared environment verification PASS

```

In [21]: # --- 7.6 Seed Reproducibility Test -----
def _short_train_eval(seed, episodes=8):
    ag = DDQNAgent(seed=seed)
    run_episodes(ag, CANEEnv(seed=1000 + seed), n_episodes=episodes, learn=True, gr
m, _, _ = run_episodes(ag, CANEEnv(seed=7000 + seed), seeds=EVAL_SEEDS[:20],
    return m["reward_mean"], ag

_r1, _a1 = _short_train_eval(0)
_r2, _a2 = _short_train_eval(0)
_r3, _a3 = _short_train_eval(1)

assert _r1 == _r2
assert all(torch.equal(p, q) for p, q in zip(_a1.online_net.parameters(), _a2.onlin

print("7.6 reproducibility PASS")

```

7.6 reproducibility PASS

8. Agent Registry & Execution Driver

Registers baselines and the Double DQN agent, then executes evaluations across all user archetypes.

8.1 Guaranteeing contact: the minimum-contact wrapper

Retuning `gamma` makes contact *affordable*; it does not make it *certain*. Whether a given seed breaks its silence is still a lottery, and a notification channel that goes quiet for a whole week on some users has stopped being a notification channel.

`MinContactAgent` removes the lottery by removing silence from the feasible set. It enforces a quota of one send per rolling 24 hours and changes nothing else:

- the **quota** decides *whether* a send happens;
- the **policy** still decides *when* within the day, and *which* message;
- `update` is forwarded untouched, so the network learns from the forced transitions too -- including that a forced send at a bad hour was expensive, and one at a good hour paid.

The reward function, the environment and the network are all unchanged. This is the standard budgeted / constrained-MDP formulation: the constraint lives on the policy's feasible set, never on its objective.

The deadline is per rolling day rather than per week on purpose. With a single deadline at the end of the episode, a silent agent would take all of its sends in the closing hours of the week, landing on whatever clock hour the episode happened to reach by then -- an artifact of the deadline rather than a choice, and it would destroy the very thing being measured, which is the hour the agent picks for this person. Per day, the agent chooses freely among 24 hours and is forced only in the last hour that can still satisfy the quota.

```
In [22]: class MinContactAgent:
          """Wrap any agent so it contacts the user at least `min_sends` per window.

          Deliberately thin: the only thing it takes away from the inner agent is the
          option of reaching the end of a window having sent nothing at all.
          """

          def __init__(self, inner, min_sends=1, window=24, label=None):
              self.inner = inner
              self.min_sends = int(min_sends)
              self.window = int(window)
              self._label = label
              self.reset()

          def reset(self):
              self.t = 0
              self.sent_in_window = 0
              if hasattr(self.inner, "reset"):
```



```

        self.inner.reset()

@property
def name(self):
    # Defaults to the plain algorithm name: `run_episodes` records
    # `agent.name` into the results table, and all four notebooks have to
    # agree on that label for the cross-algorithm comparison to line up.
    base = getattr(self.inner, "name", type(self.inner).__name__)
    return self._label or (base() if callable(base) else base)

def _quota_forces_send(self):
    """True when holding now would miss this window's quota."""
    hours_left = self.window - (self.t % self.window)
    remaining = self.min_sends - self.sent_in_window
    return remaining > 0 and hours_left <= remaining

def act(self, state, greedy=False):
    action, aux = self.inner.act(state, greedy=greedy)

    if self._quota_forces_send() and action == HOLD:
        # Forced to send. *Which* message remains the inner agent's call:
        # take its highest-valued non-Hold action rather than imposing one.
        q = aux.get("q_values") if isinstance(aux, dict) else None
        if q is not None and len(q) == N_ACTIONS:
            action = int(np.argmax(np.asarray(q[1:], dtype=float))) + 1
        else:
            action = ENGAGE
        if isinstance(aux, dict):
            aux = {**aux, "forced": True}

    self.t += 1
    self.sent_in_window += int(action != HOLD)
    if self.t % self.window == 0:
        self.sent_in_window = 0
    return action, aux

def update(self, state, action, reward, next_state, done, aux):
    out = self.inner.update(state, action, reward, next_state, done, aux)
    if done:
        self.t = 0
        self.sent_in_window = 0
    return out

def state_dict_for_save(self):
    return self.inner.state_dict_for_save()

def __getattr__(self, item):
    # Reached only for attributes this class does not define, so it cannot
    # shadow act/update/reset above. This is what lets `.history`,
    # `.online_net`, `.cfg` and the rest of the diagnostics below reach the
    # inner network unchanged.
    return getattr(self.__dict__["inner"], item)

```

MIN_SENDS, CONTACT_WINDOW = 1, 24

```

def make_ddqn(seed=0, label="DDQN"):
    """The Double DQN agent as shipped: retuned, and guaranteed to make contact.

    The single construction point for every reported result, so the learning
    curves in section 9 and the checkpoints in section 12 describe the same
    policy the leaderboard scores. Set MIN_SENDS = 0 to recover the
    unconstrained agent and reproduce the collapse.
    """
    inner = DDQNAgent(seed=seed)
    if not MIN_SENDS:
        return inner
    return MinContactAgent(inner, min_sends=MIN_SENDS,
                           window=CONTACT_WINDOW, label=label)

# Self-test: a wrapped agent left to hold forever must still send once a day.
_probe_mc = MinContactAgent(FixedScheduleAgent(hour=18), min_sends=1, window=24)
_probe_mc.inner.act = lambda state, greedy=False: (HOLD, {})
_probe_mc.reset()
_forced = sum(_probe_mc.act(None)[0] != HOLD for _ in range(168))
assert _forced == 7, f"expected 7 forced sends in a week, got {_forced}"

print(f"MinContactAgent: >={MIN_SENDS} send per {CONTACT_WINDOW}h "
      f"({_forced} forced sends a week for an always-hold policy); "
      f"make_ddqn() -> {make_ddqn().name!r}")

```

MinContactAgent: >=1 send per 24h (7 forced sends a week for an always-hold policy);
make_ddqn() -> 'DDQN'

```

In [23]: AGENTS = {}

def register(name, factory, kind="learning"):
    if name in AGENTS:
        raise ValueError(f"{name} already registered")
    AGENTS[name] = {"factory": factory, "kind": kind}
    return name

register("Fixed-18:00", lambda s: FixedScheduleAgent(hour=18), kind="baseline")
register("Random",      lambda s: RandomAgent(seed=s),      kind="baseline")
register("DDQN",        lambda s: make_ddqn(seed=s))

print(f"{len(AGENTS)} registered:", ", ".join(AGENTS))

```

3 registered: Fixed-18:00, Random, DDQN

```

In [24]: import time

# --- Evaluation sweep -----
# Every agent against every archetype. Per-archetype rather than pooled:
# averaging over the five user types hides the result, because a policy
# can be correct for one archetype and clearly wrong for another.
RESULTS_ALL = {}
_t0 = time.time()

def run_all(archetype, names=None, n_seeds=None, train_episodes=TRAIN_EPISODES):
    out = {}

```

```

for name in (names or list(AGENTS)):
    spec = AGENTS[name]
    if spec["kind"] == "baseline":
        r = evaluate_baseline(spec["factory"], archetype=archetype)
    else:
        r = train_and_evaluate(spec["factory"], n_seeds=n_seeds,
                               train_episodes=train_episodes,
                               archetype=archetype, collect_hours=True)

    r["agent"] = name
    out[name] = r
    print(f" {archetype:16} {name:12} reward {r['reward_mean']:8.2f} "
          f"ctr {r['ctr']:.3f} sends {r['sends_per_episode']:6.2f} "
          f"optout {r['optout_rate']:.3f}")

return out

for arch in ARCHETYPES:
    RESULTS_ALL[arch] = run_all(arch)
RESULTS = {a: RESULTS_ALL[a] for a in FOCUS}

print(f"\n[{time.time() - _t0:.0f}s for full evaluation sweep]")

```

00	OfficeWorker	Fixed-18:00	reward	2.01	ctr 0.316	sends 7.00	optout 0.0
95	OfficeWorker	Random	reward	-121.53	ctr 0.070	sends 29.14	optout 0.9
00	OfficeWorker	DDQN	reward	-19.31	ctr 0.230	sends 33.38	optout 0.0
00	NightOwlStudent	Fixed-18:00	reward	-18.14	ctr 0.060	sends 7.00	optout 0.0
95	NightOwlStudent	Random	reward	-132.32	ctr 0.058	sends 31.00	optout 0.9
03	NightOwlStudent	DDQN	reward	11.08	ctr 0.479	sends 22.23	optout 0.0
00	NightShiftWorker	Fixed-18:00	reward	-4.18	ctr 0.238	sends 7.00	optout 0.0
95	NightShiftWorker	Random	reward	-129.61	ctr 0.068	sends 31.14	optout 0.9
00	NightShiftWorker	DDQN	reward	-12.02	ctr 0.147	sends 8.29	optout 0.0
00	NormalStudent	Fixed-18:00	reward	-7.39	ctr 0.197	sends 7.00	optout 0.0
95	NormalStudent	Random	reward	-133.41	ctr 0.064	sends 31.89	optout 0.9
00	NormalStudent	DDQN	reward	-10.04	ctr 0.269	sends 20.16	optout 0.0
00	Housewife	Fixed-18:00	reward	-7.62	ctr 0.194	sends 7.00	optout 0.0
90	Housewife	Random	reward	-111.60	ctr 0.098	sends 29.13	optout 0.9
18	Housewife	DDQN	reward	43.95	ctr 0.327	sends 46.33	optout 0.0

[2710s for full evaluation sweep]

In [25]: `from pathlib import Path`

```

_resdir = Path("results")
_resdir.mkdir(exist_ok=True)

# --- Machine-readable export -----
# One row per (archetype, agent, seed) in the 7-column schema shared by
# all four notebooks, so the result files can be concatenated directly
# for the cross-algorithm comparison and the dashboard.
rows = []
for arch, res in RESULTS_ALL.items():
    for name, r in res.items():
        if "per_seed_rewards" in r:
            for si, sr in enumerate(r["per_seed_rewards"]):
                rows.append({"archetype": arch, "agent": name, "seed": si,
                            "reward_mean": sr, "ctr": r["ctr"],
                            "sends_per_episode": r["sends_per_episode"],
                            "optout_rate": r["optout_rate"]})
        else:
            rows.append({"archetype": arch, "agent": name, "seed": 0,
                        "reward_mean": r["reward_mean"], "ctr": r["ctr"],
                        "sends_per_episode": r["sends_per_episode"],
                        "optout_rate": r["optout_rate"]})

pd.DataFrame(rows).to_csv(_resdir / "ddqn_results.csv", index=False)
print(f"Results saved to {_resdir / 'ddqn_results.csv'}")

```

Results saved to results\ddqn_results.csv

9. Learning Curves & Optimisation Diagnostics

Visualises policy learning over time (D1) and loss/Q-value optimization metrics (D2).

```

In [26]: ddqn_trace_agent = make_ddqn(seed=0)
         _snap_result = train_and_evaluate(
             lambda s: make_ddqn(seed=s), n_seeds=1,
             snapshot_every=50, snapshot_seed=0)

         ddqn_trace_agent = make_ddqn(seed=0)
         _trace_env = CANEEEnv(seed=1000)
         run_episodes(ddqn_trace_agent, _trace_env, n_episodes=TRAIN_EPISODES, learn=True, g

```

```
Out[26]: ({'agent': 'DDQN',
  'reward_mean': -135.56070862312694,
  'reward_std': 79.62779895774766,
  'ctr': 0.10446455901001356,
  'sends_per_episode': 38.115,
  'clicks_per_episode': 3.9816666666666665,
  'optout_rate': 0.64,
  'episode_length': 100.21666666666667},
None,
array([-264.6947878 , -112.80446773, -103.70527866, -26.2267764 ,
        -59.57430801, -152.58612762, -202.2579351 , -135.98708121,
        -45.23501257, -77.74225208, -214.83960519, -92.41890996,
        -150.0851078 , -295.90499744, -112.2765904 , -141.77764677,
        -135.54812621, -171.21974118, -46.44516536, -89.95437356,
        -100.29945495, -73.83531803, -52.64084068, -32.06777626,
        -23.31484454, -45.21878111, -217.59251388, -67.52371603,
        -222.75086274, -104.12478464, -53.5936656 , -124.79994571,
        -38.24513676, -298.17547377, -62.98141681, -154.12892941,
        -97.19204857, -184.80112324, -76.36177446, -71.42952852,
        -63.58906583, -213.89393484, -166.26876664, -66.82204187,
        -96.06785719, -130.34700215, -75.53055491, -142.15592876,
        -62.91722265, -196.90379758, -59.49696365, -121.20866373,
        -173.718053 , -165.69637955, -47.31625952, -60.82486612,
        -72.05583217, -197.94588494, -54.81603286, -135.44456923,
        -71.86145933, -44.91310328, -43.99662327, -44.15050932,
        -59.22241745, -161.81331667, -124.28697488, -83.14280604,
        -380.52337916, -102.49584046, -67.22818194, -163.40619059,
        -81.54518631, -39.15144072, -134.91653984, -42.57955975,
        -194.52208823, -140.68211731, -26.97639435, -80.95297113,
        -50.4509973 , -79.9366276 , -45.72340277, -328.41530276,
        -41.99613856, -241.08781429, -187.50998369, -71.59765915,
        -122.89041785, -214.54246211, -80.40578337, -66.06299032,
        -52.59151562, -64.47664516, -133.90088636, -290.16225699,
        -111.23667406, -329.90252464, -99.35072173, -62.85066334,
        -232.44577169, -86.47918718, -157.48711863, -241.81436073,
        -27.20846968, -334.72171152, -97.36899394, -241.33384412,
        -169.13159881, -129.25958237, -56.05530545, -77.03800168,
        -28.08977883, -117.58788748, -110.61358648, -122.82857406,
        -55.321644 , -54.24333159, -75.30920055, -272.30376225,
        -171.37580316, -71.03585489, -89.25908727, -106.74315226,
        -65.92426593, -121.40489862, -69.04818458, -60.20788693,
        -151.60686613, -337.15688936, -231.74158362, -213.56731967,
        -103.48690017, -136.17118773, -44.29787714, -164.78683423,
        -265.24942783, -244.71176359, -41.62248185, -85.56611312,
        -176.00988905, -43.16056612, -125.37839852, -286.86231668,
        -158.56542455, -265.64909163, -33.89486525, -80.91478928,
        -143.70731814, -159.52642232, -52.90183666, -57.72638384,
        -59.6434157 , -438.10296548, -65.53481035, -108.65204039,
        -150.92721644, -342.61648175, -225.34395305, -90.59026078,
        -158.23328766, -184.56522189, -180.96814084, -96.11467478,
        -85.16387869, -84.25966718, -277.95716103, -61.00299218,
        -260.61011113, -259.16540069, -133.53199128, -85.86581443,
        -89.46086792, -123.499275 , -367.83546678, -200.65617799,
        -50.98674265, -148.59991427, -70.84586071, -349.74459155,
        -47.31138337, -90.34903655, -33.00979769, -266.13942275,
        -218.74954791, -120.50886981, -258.34828794, -179.54621087,
```

-98.95925581, -74.78272391, -100.12087447, -82.27914938,
-131.89896366, -75.65343687, -94.70710708, -59.50984507,
-240.38233322, -350.72529073, -90.41318659, -96.43667944,
-115.0989404 , -101.86495321, -49.64854231, -101.84303081,
-113.83026543, -253.85169826, -162.41497613, -41.11788604,
-57.81261203, -61.7030791 , -67.06814016, -149.03263437,
-42.77386821, -43.41331283, -174.82541173, -197.57403826,
-250.46286573, -320.15140003, -190.12649369, -75.18002208,
-123.45310028, -151.64850691, -146.08777811, -115.50977005,
-236.56987821, -258.07501133, -148.20136996, -51.53297287,
-223.88000361, -237.92277412, -184.87691977, -192.37295823,
-65.94239481, -64.81052676, -200.63715326, -151.79030194,
-173.57044848, -28.27861901, -284.02724165, -185.01516821,
-236.4220798 , -67.22007801, -127.47225679, -233.1750289 ,
-109.62168996, -61.44409674, -185.44266932, -351.13610636,
-161.9864025 , -191.30376644, -98.27625483, -50.10245248,
-61.76639916, -120.36989194, -97.64251259, -273.6176103 ,
-186.04526872, -323.99391808, -151.24530802, -48.1740756 ,
-132.05485294, -87.17472001, -150.62542926, -244.01513628,
-91.24217166, -281.51654342, -179.94621063, -252.27828156,
-158.8313044 , -84.79413135, -205.58495458, -324.56051499,
-112.56436319, -75.25729929, -191.68737248, -90.92442066,
-312.35389942, -110.71321679, -102.34670628, -340.43456365,
-338.54170383, -228.95044578, -309.95118388, -105.71365046,
-249.65854519, -121.44453258, -42.88188327, -71.50069723,
-295.95888309, -215.98444578, -233.61083559, -268.1521871 ,
-97.97608317, -157.80922219, -85.9386942 , -140.01203957,
-80.00852536, -99.9309096 , -192.77124184, -143.7614441 ,
-77.37609711, -113.01626527, -168.31848851, -106.77783867,
-125.57581525, -78.78312627, -179.07343321, -211.27266938,
-133.66165228, -138.20287545, -278.68966892, -153.96939627,
-81.38613942, -60.63705452, -149.40563738, -51.92018651,
-85.59962969, -48.01551276, -185.13506846, -287.38463465,
-242.83031977, -129.56645949, -164.42040875, -173.94878031,
-247.91914779, -267.29989831, -53.7837014 , -271.67239622,
-267.66642372, -191.63018018, -74.59021702, -268.10451181,
-274.29788663, -110.59099254, -42.06970487, -210.13168305,
-178.15673767, -296.624358 , -196.51402002, -62.40881677,
-30.93103644, -96.41955713, -194.46587991, -89.60307149,
-104.94026988, -230.43708435, -243.14008834, -248.97744847,
-57.38620015, -152.02242371, -221.62128527, -275.53075531,
-232.81968385, -237.00497753, -232.2702409 , -264.8273611 ,
-252.15925834, -57.34032273, -94.53042951, -77.03837744,
-38.42764508, -257.23086099, -34.74640292, -262.22148603,
-255.22076064, -250.68738099, -98.17945178, -169.68945341,
-192.9397702 , -138.79335796, -223.62766858, -84.79651817,
-54.81295834, -129.13043305, -232.83218207, -184.9949613 ,
-254.42667465, -267.25893751, -92.2048062 , -262.88474014,
-202.26613239, -124.11262073, -191.43358191, -261.83053638,
-124.74803579, -179.76909637, -228.98052931, -36.70448274,
-165.04452428, -139.96395778, -194.8609506 , -99.6852118 ,
-228.21567538, -164.48159604, -254.57915296, -68.4158985 ,
-191.69682374, -184.52063166, -84.12097239, -127.13136781,
-212.35224179, -257.26526574, -175.77306473, -241.36593708,
-216.67126221, -232.37486469, -125.64743536, -44.56727668,
-218.86763246, -186.1046579 , -173.68146633, -209.36645578,

```

-114.2145589 , -40.69473981, -174.01494585, -210.71704778,
-225.68673635, -37.04287512, -187.34856786, -194.28548512,
-193.21280328, -94.84379498, -242.14100302, -135.92892144,
-99.25137793, -111.48691371, -158.60403969, -172.5779382 ,
-250.78748345, -223.05400361, -91.86919796, -128.05837271,
-104.61818386, -79.49530632, -217.84610135, -205.71987643,
-205.5327624 , -282.0511652 , -160.80474064, -221.85716122,
-207.82480805, -231.75477373, -212.4271025 , -69.78435667,
-136.04830338, -152.02847356, -156.41549611, -144.78284595,
-155.04832349, -161.12081691, -113.47655704, -181.76788157,
-158.27104332, -106.50728419, -43.79235578, -187.4444038 ,
-196.14313992, -200.81399187, -147.45701205, -113.67493453,
-96.0366842 , -168.78996633, -161.18021961, -176.65025194,
-188.36471488, -109.45958075, -72.48759325, -76.23704379,
-124.53049408, -148.80143848, -175.75295946, -214.44554048,
-105.89980544, -94.9370593 , -146.92786743, -142.79351537,
-161.43081009, -190.49791284, -150.12207843, -157.67957502,
-186.99506103, -187.57606782, -194.93640535, -121.39934991,
-140.80302223, -143.61628897, -132.40978544, -119.903556 ,
-84.83308808, -150.05495974, -178.97767988, -94.3518634 ,
-178.95179116, -63.39900196, -79.04676015, -88.47980368,
-193.45411885, -178.14382831, -108.29668171, -97.09118965,
-129.66479078, -76.95961576, -107.88368202, -100.19681072,
-196.54496493, -47.80233361, -108.13287855, -87.72784964,
-59.35738233, -118.47573813, -86.8799641 , -90.34755938,
-85.39549831, -140.87014689, -78.75418249, -110.55130742,
-179.7830093 , -53.85810834, -70.83647668, -141.6612318 ,
-73.72248313, -136.6579812 , 2.3292328 , -93.28569816,
-162.07936459, -61.48028462, -61.68355389, -67.13559937,
-90.44713573, -52.52193406, -86.96680393, -66.70379828,
-124.55917709, -99.75072193, -132.65807243, -117.33897878,
-80.54825036, -80.50926989, -139.69654164, -70.22631977,
-68.84025173, -60.95371845, -52.2514158 , -50.75553902,
-67.96423396, -80.39179664, -53.31878957, -96.31131194,
-90.77175776, -70.76812101, -67.03105454, -112.42442536,
-100.41510113, -82.83514954, -56.14876204, -70.98348129,
-30.48299493, -74.05846238, -53.25433029, -59.36565421,
-40.13153853, -21.21789922, -56.71453115, -55.75739912,
-36.49639677, -56.19352467, -65.82244013, -59.86795727,
21.61384969, -72.37095178, -83.21742021, -12.7119608 ,
-65.82592527, -99.3264844 , -42.60075553, -61.42106816,
-27.35179305, -96.31266962, -35.86613785, -17.60799498,
-48.64613643, -50.84066491, -76.72658593, -10.55829285,
-10.82329421, -69.62670931, -66.33954589, -47.02350941,
-39.24302055, -49.76105128, -63.119449 , -67.66846822,
-0.70612484, -66.18996012, -58.3213775 , -51.784974 ,
-16.8037437 , 26.17720452, 47.80041996, -29.71964385]],
array([[1627, 441, 451],
[1583, 498, 434],
[1579, 457, 484],
[1625, 442, 452],
[1558, 479, 477],
[1591, 475, 439],
[1591, 473, 443],
[1596, 460, 441],
[1585, 464, 452],

```

```


[1624, 448, 426],
[1559, 489, 445],
[1623, 446, 425],
[1550, 486, 466],
[1557, 496, 452],
[1485, 553, 474],
[1523, 535, 449],
[1566, 518, 421],
[1544, 538, 421],
[1511, 522, 470],
[1462, 550, 486],
[1481, 556, 463],
[1436, 575, 489],
[1490, 559, 453],
[1515, 519, 477]])

```

```

In [27]: # --- Plot D1: Learning Curve ---
snaps = _snap_result.get("snapshots", [])
if snaps:
    fig, axes = plt.subplots(1, len(FOCUS), figsize=(5.5 * len(FOCUS), 4.0), sharey=True)
    for col, arch in enumerate(FOCUS):
        ax = axes[0, col]
        eps_list = [s[0] for s in snaps]
        rew_list = [s[1][arch]["reward"] for s in snaps]
        ax.plot(eps_list, rew_list, "o-", color=ARCH_COLOUR[arch], lw=1.8, markersize=10)
        ax.axhline(0, color="#999999", lw=0.8, ls=":")
        ax.set_xlabel("training episodes")
        ax.set_ylabel("greedy reward / week" if col == 0 else "")
        ax.set_title(arch, fontsize=10)
        ax.grid(alpha=0.3)
    fig.suptitle("D1 DDQN Learning Curve", fontsize=11)
    fig.tight_layout()
    save(fig, "D1_ddqn_learning_curve")
    plt.show()

```



```

In [28]: # --- Plot D2: Optimiser Diagnostics ---
hist = pd.DataFrame(ddqn_trace_agent.history)

if len(hist) > 0:
    fig, axes = plt.subplots(1, 4, figsize=(14, 3.0))
    for ax, (col, lab) in zip(axes, [("loss", "loss (MSE)"), ("mean_q", "mean Q"),
                                     ("mean_v", "mean V"), ("mean_r", "mean R")]):
        ax.plot(hist["update"], hist[col], lw=1.0, color="#2f6fb5", alpha=0.6)


```



```

if len(hist) > 20:
    smoothed = hist[col].rolling(window=min(50, len(hist)//5), min_periods=
    ax.plot(hist["update"], smoothed, lw=2.0, color="#c4453c")
ax.set_xlabel("update")
ax.set_title(lab, fontsize=9)
ax.grid(alpha=0.3)
fig.suptitle("D2 DDQN Optimiser Diagnostics", fontsize=11)
fig.tight_layout()
save(fig, "D2_ddqn_diagnostics")
plt.show()

```



10. Decision Distribution & Policy Behaviour

Plots summary rewards by archetype (D4) and hourly decision breakdown vs ground-truth policy (D3).

```

In [29]: # --- Plot D4: Reward by Archetype ---
names = list(AGENTS)
x = np.arange(len(ARCHETYPES))
width = 0.8 / len(names)

fig, ax = plt.subplots(figsize=(11, 3.6))
for i, name in enumerate(names):
    vals = [RESULTS_ALL[a][name]["reward_mean"] for a in ARCHETYPES]
    errs = [RESULTS_ALL[a][name]["reward_sem_std"] for a in ARCHETYPES]
    ax.bar(x + i * width - 0.4 + width / 2, vals, width * 0.92, yerr=errs,
          label=name, error_kw={"lw": 0.8, "ecolor": "#333333"})
ax.axhline(0.0, color="#333333", lw=0.8)
ax.set_xticks(x)
ax.set_xticklabels(ARCHETYPES, fontsize=8)
ax.set_ylabel("greedy reward / week")
ax.legend(frameon=False, fontsize=8, ncol=len(names))
ax.set_title("D4 Reward by Agent and Archetype (+/- 1 SD)", fontsize=11)
fig.tight_layout()
save(fig, "D4_ddqn_reward_by_archetype")
plt.show()

```


D3 DDQN Decision Distribution by Hour



```
In [31]: # --- Plot D5: Hyperparameter Sensitivity ---
_search_rng = np.random.default_rng(2026)
_N_SEARCH = 20

_search_results = []
for i in range(_N_SEARCH):
    _lr = float(10 ** _search_rng.uniform(-5, -3))
    _eps_end = float(_search_rng.uniform(0.01, 0.15))
    _eps_decay = int(_search_rng.choice([20000, 50000, 80000, 100000]))
    _target_freq = int(_search_rng.choice([200, 500, 1000, 2000]))
    _buf_size = int(_search_rng.choice([5000, 10000, 20000]))
    _hidden = tuple(_search_rng.choice([(32, 32), (64, 64), (128, 128)]))

    try:
        r = train_and_evaluate(
            lambda s, _lr=_lr, _eps_end=_eps_end, _eps_decay=_eps_decay,
                _target_freq=_target_freq, _buf_size=_buf_size,
                _hidden=_hidden:
                DDQNAgent(seed=s, lr=_lr, epsilon_end=_eps_end,
                    epsilon_decay_steps=_eps_decay,
```

```

        target_update_freq=_target_freq,
        buffer_size=_buf_size,
        hidden_sizes=_hidden),

    n_seeds=2,
    archetype="OfficeWorker")
_search_results.append({
    "lr": _lr, "eps_end": _eps_end, "eps_decay": _eps_decay,
    "target_freq": _target_freq, "buf_size": _buf_size,
    "hidden": str(_hidden), "reward": r["reward_mean"],
    "ctr": r["ctr"], "sends": r["sends_per_episode"],
    "optout": r["optout_rate"]})
except Exception as ex:
    print(f"FAILED: {ex}")


_search_df = pd.DataFrame(_search_results)

if len(_search_df) > 0:
    _hp_cols = ["lr", "eps_end", "eps_decay", "target_freq", "buf_size"]
    fig, axes = plt.subplots(1, len(_hp_cols), figsize=(3.2 * len(_hp_cols), 3.0))

    for ax, col in zip(axes, _hp_cols):
        ax.scatter(_search_df[col], _search_df["reward"], s=30, alpha=0.7,
                    color="#2f6fb5", edgecolors="#333333", linewidth=0.5)
        ax.axhline(0, color="#999999", lw=0.8, ls=":")
        ax.set_xlabel(col, fontsize=8)
        ax.set_ylabel("reward" if col == _hp_cols[0] else "")
        ax.grid(alpha=0.3)
        if col == "lr":
            ax.set_xscale("log")

    fig.suptitle("D5 DDQN Random Search – OfficeWorker", fontsize=11)
    fig.tight_layout()
    save(fig, "D5_ddqn_random_search")
    plt.show()

```



12. Statistical Robustness & Model Persistence

Plots per-seed reward variance (D6) and saves trained model weights.

```

In [32]: # --- Plot D6: Robustness Boxplot ---
fig, axes = plt.subplots(1, len(FOCUS), figsize=(5.5 * len(FOCUS), 4.0), sharey=True)

for col, arch in enumerate(FOCUS):
    ax = axes[0, col]
    agent_names = []

```

```


agent_rewards = []
for name in AGENTS:
    r = RESULTS_ALL[arch][name]
    if "per_seed_rewards" in r and len(r["per_seed_rewards"]) > 1:
        agent_names.append(name)
        agent_rewards.append(r["per_seed_rewards"])

if agent_rewards:
    positions = range(len(agent_names))
    bp = ax.boxplot(agent_rewards, positions=list(positions), widths=0.5, patch
for patch in bp["boxes"]:
    patch.set_facecolor("#2f6fb5")
    patch.set_alpha(0.5)
ax.set_xticks(list(positions))
ax.set_xticklabels(agent_names, fontsize=8)

ax.axhline(0, color="#999999", lw=0.8, ls=":")
ax.set_ylabel("greedy reward / week" if col == 0 else "")
ax.set_title(arch, fontsize=10)
ax.grid(alpha=0.3, axis="y")

fig.suptitle("D6 DDQN Robustness — Per-Seed Reward Distribution", fontsize=11)
fig.tight_layout()
save(fig, "D6_ddqn_robustness")
plt.show()

```



```

In [33]: # --- Model Checkpointing ---
_modeldir = Path("models")
_modeldir.mkdir(exist_ok=True)

for seed in range(N_SEEDS):
    agent = make_ddqn(seed=seed)
    train_env = CANEnv(seed=1000 + seed)
    run_episodes(agent, train_env, n_episodes=TRAIN_EPISODES, learn=True, greedy=False)
    torch.save(agent.state_dict_for_save(), _modeldir / f"ddqn_seed{seed}.pt")
    print(f"Saved model: ddqn_seed{seed}.pt")

print("All models successfully saved.")

```

Saved model: ddqn_seed0.pt
Saved model: ddqn_seed1.pt
Saved model: ddqn_seed2.pt
Saved model: ddqn_seed3.pt
Saved model: ddqn_seed4.pt
All models successfully saved.

Summary & Section Breakdown

This notebook is structured into 12 self-contained sections:

1. **Setup & Dependencies:** Environment configuration, imports, action constants, and feature index mapping.
2. **Agent Interface:** `Agent` abstract base class establishing the contract for all policies.
3. **Environment Dynamics:** `CANEEEnv` implementation modeling receptive windows, fatigue accumulation, click probability, and churn hazard.
4. **Baselines & Evaluation Protocol:** Baseline policies (`FixedScheduleAgent` , `RandomAgent`) and evaluation harness (`run_episodes` , `train_and_evaluate` , `EVAL_SEEDS`).
5. **Sanity Checks & Environment EDA:** Unit verification of fatigue mechanics and initial environment visualization.
6. **Double DQN Agent:** Core algorithm implementation containing `DDQN_CONFIG` , `ReplayBuffer` , `QNetwork` , and target-decoupled `DDQNAgent` .
7. **Correctness Verification Tests:** Unit test suite (7.1 - 7.6) confirming interface compliance, target decoupling, evaluation purity, synthetic problem convergence, and seed reproducibility.
8. **Registry & Execution Driver:** Multi-seed benchmark execution driver and CSV results exporter (`ddqn_results.csv`).
9. **Learning Curves & Diagnostics:** Training progress curve (D1) and optimization metrics trace (D2: loss, Q-values, epsilon).
10. **Decision Distribution & Policy Behaviour:** Reward comparisons (D4) and 24-hour action distribution vs. receptive windows (D3).
11. **Hyperparameter Sensitivity Analysis:** Random search over hyperparameters (D5).
12. **Statistical Robustness & Model Persistence:** Seed variance boxplots (D6) and PyTorch checkpoint saving (`ddqn_seed{0..4}.pt`).

Figures Summary

- **D1 (Learning Curve):** Evaluates greedy weekly reward across probe training snapshots.
- **D2 (Optimiser Diagnostics):** Tracks MSE loss, mean Q -value, max Q -value, and linear ϵ -decay per optimization update.
- **D3 (Decision Distribution):** Hourly stacked bar charts of policy action selection vs. baseline responsiveness and ground-truth policy.

- **D4 (Reward by Archetype):** Grouped bar chart comparing weekly greedy reward across Fixed-18:00, Random, and Double DQN.
- **D5 (Random Search Scatter):** Analyzes reward sensitivity across learning rate, ϵ -decay, target update frequency, and buffer capacity.
- **D6 (Robustness Boxplot):** Displays reward variance across random seeds to verify statistical consistency.

13. Key findings and implications

Results

Averaged over five training seeds and the 200 held-out evaluation episodes:

Archetype	Reward	CTR	Sends/week	Opt-out
Housewife	+43.96	0.327	46.3	0.018
NightOwlStudent	+11.08	0.479	22.2	0.003
NightShiftWorker	-12.02	0.147	8.3	0.000
NormalStudent	-10.04	0.269	20.2	0.000
OfficeWorker	-19.31	0.230	33.4	0.000
Mean	+2.73	0.290	26.1	0.004

Reference points: Fixed-18:00 averages **-7.06**, Random **-125.69**, DQN **+6.55**. This run uses the retuned configuration of section 1 together with the minimum-contact wrapper of section 8.1.

Observations and patterns

1. **Coverage went from 2 of 5 archetypes to 5 of 5.** Under the original configuration ($\gamma = 0.99$, no wrapper) this agent contacted only Housewife and NightOwlStudent and scored exactly 0.00 on the other three. Lowering γ to 0.90 and adding the quota moved it to full coverage.
2. **Double DQN does not beat DQN here (+2.73 against +6.55).** The decoupled target that defines the algorithm is designed to reduce over-estimation bias, and over-estimation was never the failure mode in this environment.
3. **The whole of the gap is on NightShiftWorker:** -12.02 here against DQN's -0.23. DQN learned to leave that user alone; the wrapper forbids this agent from doing so and obliges it to send about 8 times a week at a CTR of 0.147, less than half of break-even.

Key findings and their implications

Finding 1 — the never-send collapse was an economic result, not an optimisation failure. The agent was not stuck, and it was not under-trained. At $\gamma = 0.99$ across a 168-step episode, the discounted fatigue tail of a single send outweighs a click that pays once, so silence was the correct answer to the question actually being asked. Changing γ changed the question. This is the single most important finding in this notebook, and it generalises: **when an RL agent does nothing, check the reward geometry before the optimiser.**

Finding 2 — the cost of guaranteed coverage is measurable, and it is about 6 reward per user forced. The minimum-contact wrapper is a constraint on the policy's *feasible set*, never on its objective -- the standard budgeted-MDP formulation. It buys a product guarantee (no user is ever abandoned for a whole day) and it pays for that guarantee in reward on exactly the users where contact is unprofitable. Both halves of that trade belong in the report.

Finding 3 — algorithmic sophistication did not substitute for correct problem specification. Double DQN is the more advanced estimator and it finished second to plain DQN. The variable that decided the outcome was γ , which both algorithms share.

Implication for deployment. Guaranteed contact and maximum reward are genuinely different objectives, and the choice between them is a product decision rather than a modelling one. `MIN_SENDS = 0` in section 8.1 recovers the unconstrained agent and reproduces the collapse, so both policies remain available and the trade can be re-measured at any time.

How to read a reward of `0.00`

`0.00` is not a missing value and not a crash. It is exactly what an agent earns by never sending: no clicks are collected, so no click reward is paid, and no send or fatigue cost is incurred either. Because three of the five archetypes score *negatively* under the Fixed-18:00 baseline, silence genuinely outscores the industry-standard schedule on those users. Any agent that discovers this is not malfunctioning -- it has found a real, and uncomfortable, local optimum.

The economics that govern every result here

A send is worth making only when its click probability clears the break-even click-through rate

$$\text{CTR}_{\text{break-even}} = \frac{W_{\text{send}} + W_{\text{fat}} \cdot \kappa / (1 - \lambda)}{R_{\text{click}}} = 0.340$$

The $\kappa / (1 - \lambda)$ term is the whole story: because fatigue decays geometrically rather than instantly, one send commits the agent to a stream of future fatigue penalties. At $\gamma = 0.99$ over a 168-step episode that discounted tail is worth roughly a hundred steps of debt while

the click pays once, which is why silence can dominate. Every finding above is a consequence of that single inequality.